

The Revolutionary Beginner's Guide to Quantum Computing

Why We Don't Have Quantum Computers Yet — and What the Geometric Alternative Offers

2026-08-21

Abstract

A 20-chapter beginner's guide to quantum computing that starts with an honest question — why don't we have quantum computers yet? — rather than the usual “here's what a qubit is.” The guide teaches the standard curriculum (qubits, entanglement, algorithms, error correction) accurately, explains why the standard approach faces a thermodynamic wall (a 20,000-fold cooling gap that is Carnot-limited), and introduces the ultrametric alternative: passive geometric fault tolerance on Bruhat-Tits trees, with DOI-registered evidence, computationally validated error thresholds 75 times higher than surface codes (validated by classical simulation; not yet demonstrated on quantum hardware), and three falsifiable predictions that can be tested for under \$300,000.

Author: [Rowan Brad Quni-Gudzin](#) **ORCID:** [0009-0002-4317-5604](#) **DOI:** [10.5281/zenodo.22040750](#)

Part I: The Honest Landscape

How to Read This Guide

This is not a standard quantum computing textbook. Most guides start with a qubit, walk you through superposition and entanglement with colorful Bloch sphere diagrams, present a few algorithms, and end with an optimistic note about the quantum future.

This guide is different — and it is honest about how. **It is an opinionated textbook.** It teaches the standard quantum computing curriculum accurately and fairly (Chapters 4–7), because you need to know that material. But it also argues for a specific thesis — that the standard approach faces a thermodynamic wall, and that a geometric alternative deserves serious attention. The author is the originator of that alternative and is transparent about it: the work being advocated for is cited throughout via DOI-registered publications. You do not need to agree with the thesis to learn from the guide. But you do need to know that this is not a neutral survey — it is an argued case, and every claim is tagged with its confidence level so you can evaluate it yourself.

This guide starts with a different question: **Why don't we have quantum computers yet?**

If you have been following technology news for the past decade, you have probably heard that practical quantum computers are “five to ten years away.” You may have heard this in 2015. You may have heard it in 2020. You are hearing it now, in 2026. Each year, the timeline slides forward. Each year, billions of dollars are invested. Each year, press releases announce breakthroughs. And yet — no commercially useful quantum computation has ever been performed. Not once.

This guide explains why. It introduces the standard quantum computing story — qubits, gates, entanglement, algorithms — because you need to know that story. But it also introduces the part that most guides leave out: the thermodynamic wall, the surface code plateau, and the geometric alternative that may matter more than any particular qubit technology.

Prerequisites

This guide is written for readers with a STEM background — comfort with algebra, basic probability, scientific notation, and ideally some exposure to linear algebra (vectors and matrices). If you have taken an undergraduate physics or engineering course, you have the mathematical tools you need. Concepts like the Carnot efficiency, tensor products, and the stabilizer formalism are explained as they appear, but the explanations assume mathematical maturity.

If you are a complete beginner to STEM, Chapters 4–7 (the standard quantum computing curriculum) may be challenging. Read them slowly. The Bloch sphere and quantum gates require comfort with complex numbers and matrix multiplication. If these are unfamiliar, you may want to consult a linear algebra primer alongside this guide.

That said: **the most important ideas in this guide do not require mathematics.** Chapters 1–3 (why we don't have quantum computers), Chapter 10 (what quantum delivers today), and Chapters 15–17 (how to test scientific claims) are accessible to any educated adult. You can skip the mathematical sections and still understand the core argument.

Key factual claims in this guide carry a **confidence tag** — [EST] for established fact, [PROP] for a framework proposal, [GAP] for an acknowledged gap. Not every sentence is tagged (that would make the text unreadable), but every claim that matters to the argument — especially claims about what has been demonstrated vs. what is proposed — carries a visible confidence marker. You will learn to read these tags as naturally as you read punctuation. They are the most important thing this guide teaches: how to distinguish what we know from what we hope.

Chapter 1: Why Don't We Have Quantum Computers Yet?

The Forty-Year Promise

Quantum computing was proposed in the early 1980s by physicists — most famously Richard Feynman — who recognized that simulating quantum systems on classical computers is exponentially hard. A quantum system of n particles requires 2^n numbers to describe its state. For $n = 50$, that is roughly a quadrillion numbers — more than any classical computer can store. Feynman's insight was beautifully simple: if nature is quantum mechanical, then a computer made of quantum mechanical parts should be able to simulate quantum systems efficiently. [EST]

The idea sat in physics departments for a decade. Then, in 1994, Peter Shor — a mathematician at Bell Labs — discovered something that changed everything. He showed that a quantum computer could factor large numbers exponentially faster than any known classical algorithm. [EST] Factoring numbers is the mathematical operation that underlies RSA encryption, which secures most of the internet. Suddenly, quantum computing was not just a physics curiosity. It was a threat to global digital infrastructure.

The race was on. Governments and corporations invested billions. Physicists built the first quantum logic gates using trapped ions in 1995. [EST] Superconducting qubits followed. By the 2010s, small quantum processors with a handful of qubits were operating in laboratories. In 2019, Google claimed “quantum supremacy” — a demonstration that their 53-qubit Sycamore processor could perform a specific, contrived calculation faster than the world's largest supercomputer. [EST]

And yet. Here we are in 2026. No commercially useful gate-model quantum computation has been performed by a general-purpose, programmable quantum computer. (Quantum annealing — a different approach using specialized hardware — has recently demonstrated advantage on real optimization problems; see Chapter 10. But the universal, error-corrected quantum computers that dominate industry roadmaps and media coverage have delivered nothing of commercial value.) No cryptographic key has been factored. No drug molecule has been simulated. No optimization problem has been solved faster than a classical computer could solve it. What happened?

What “Quantum Computer” Actually Means in 2026

The machines that exist today are real. They are not simulations. They perform genuine quantum gate operations on physical qubits with fidelities exceeding 99.9% on some platforms. [EST] They have demonstrated quantum error correction below the surface-code threshold — a milestone the field pursued for over twenty years. [EST] They are genuine achievements of experimental physics.

But they are not “computers” in any useful sense. They are experimental physics apparatus. They occupy specialized laboratories. Superconducting qubits require dilution refrigerators operating at 10-20 millikelvin — colder than interstellar space. Trapped ions require ultra-high-vacuum chambers with precision laser systems. These machines are hand-built, one-off devices that cost tens to hundreds of millions of dollars and require teams of PhD physicists to operate. [EST]

The gap between “this machine performed a quantum gate operation” and “this machine solved a problem I care about” is not small. It is vast. And closing it requires solving the hardest problem in the field: error correction.

The NISQ Era — What It Is, What It Isn’t

We are in what John Preskill christened the **NISQ era** — Noisy Intermediate-Scale Quantum computing. [EST] “Noisy” means the qubits are error-prone. “Intermediate-scale” means we have tens to hundreds of qubits, not millions. NISQ devices are too small and too error-prone to run the algorithms — like Shor’s — that would give quantum computers their transformative power.

The important thing to understand about the NISQ era is that **it has lasted longer than expected, and it is not ending soon**. The most optimistic roadmaps place fault-tolerant quantum computing — the kind with enough error correction to run useful algorithms — in the 2030s. More realistic assessments place it in the late 2030s or beyond. [PROP]

This is not pessimism. It is physics. The next two chapters will explain why.

How This Guide Handles Uncertainty

Before we go further, you need to understand the **confidence tags** that appear throughout this guide. Every claim carries one of five labels:

Tag	Meaning	Example
[EST]	Established — confirmed by experiment or theorem	“The strong nuclear force binds protons and neutrons.”
[PROP]	Proposed — logically consistent, not yet experimentally verified	“Tree codes achieve a depolarizing threshold of 75%.”
[GAP]	Gap — acknowledged missing step	“Full quantum validation remains future work.”
[SPEC]	Speculative — extension beyond current framework	“Application to room-temperature superconductivity.”
[OPEN]	Open question — framework cannot yet answer	“Perfect tensor existence for $p > 2$.”

These tags are not decorative. They are the most important thing this guide teaches. In a field where corporate press releases routinely blur the line between demonstrated fact and aspirational roadmap, the ability to distinguish [EST] from [PROP] is the single most valuable skill you can develop.

If you take nothing else from this guide, take this: **every time someone tells you what quantum computers will do, ask what they have actually demonstrated**. The gap between those two answers is where the real story lives.

Chapter 2: The Six Platforms — And Why None Has Won

The Basic Idea (Shared by All Platforms)

Every quantum computer is built from **qubits** — physical systems that can exist in superpositions of two states, conventionally labeled $|0\rangle$ and $|1\rangle$. To perform computation, you need three capabilities: initialize qubits to a known state, perform quantum logic gates (operations that change qubit states), and read out the results. Everything else — algorithms, error correction, compilers — is software built on top of these physical operations. [EST]

What distinguishes different quantum computing platforms is **how** they implement qubits and gates. The choice of physical system determines everything: gate speed, error rate, operating temperature, scalability, and cost. There are six major platforms competing today. None has won. All face fundamental challenges.

Superconducting Qubits — The Incumbent

Superconducting qubits are tiny electrical circuits made from superconducting materials, fabricated on chips using techniques similar to those that produce classical computer processors. They operate at microwave frequencies (2-8 GHz) and require cooling to approximately 10-20 millikelvin. [EST]

Why they lead: Superconducting qubits are the most mature platform. IBM and Google have the largest devices, with over 100 physical qubits. The software ecosystem (IBM's Qiskit, Google's Cirq) is the most developed. Google's 2024-2025 demonstration of error correction below the surface-code threshold — a milestone the field pursued for two decades — was achieved on superconducting qubits. [EST]

Why they stall: Superconducting qubits face hard physical limits. Each qubit is relatively large (millimeter scale). Each requires multiple control wires. As you add qubits, the wiring problem becomes severe — there is only so much physical space for microwave cabling entering a dilution refrigerator. The materials themselves contain two-level system (TLS) defects — microscopic imperfections that cause decoherence. Coherence times are limited to roughly a millisecond by material losses. [EST]

The fundamental issue is cooling. Dilution refrigerators at 10-20 mK provide approximately 50 microwatts of cooling power. Every qubit, every control line, every measurement adds heat. At some qubit count — far below what is needed for useful computation — the heat exceeds the cooling budget. This is not an engineering problem that better refrigerators will solve. It is a consequence of the Carnot efficiency of heat pumps operating across a temperature ratio of roughly 10^4 . [EST]

Trapped Ions — The Precision Champion

Trapped-ion quantum computers use individual atoms — typically barium, calcium, or ytterbium — stripped of one electron and suspended in a vacuum by electric fields. Qubit states are encoded in two internal atomic energy levels. Lasers manipulate these states to perform gates. [EST]

Why they excel: Trapped ions have the highest gate fidelities of any platform — exceeding 99.9%. They have the longest coherence times — up to a minute for hyperfine ground-state qubits, with the T1 lifetime effectively infinite. They are fully connected: any ion can interact with any other, a property that enables efficient implementation of certain error correction codes. [EST]

In 2024, Quantinuum became the first company to demonstrate **real-time** quantum error correction — actually closing the feedback loop rather than correcting errors in post-processing. [EST] This is a genuine milestone.

Why they stall: Trapped-ion gates are slow — typically 10-100 microseconds, compared to 10-100 nanoseconds for superconducting qubits. While proof-of-concept experiments have demonstrated nanosecond and even picosecond gates, “there is currently little research on increasing gate speeds,” according to Cornelius Hempel’s 2026 textbook chapter. [EST]

Scaling is the harder problem. A single ion trap can hold at most about 50 ions before collective motional modes become unmanageable. Two competing scaling paths exist: the quantum charge-coupled device (QCCD) architecture pursued by Quantinuum, where ions are shuttled between trap zones, and the optically linked modular approach pursued by IonQ, where separate traps are connected by photonic links. Hempel writes: “it is unclear which one ultimately scales faster and exhibits lower overhead.” [EST]

Most importantly, the trapped-ion community’s own roadmap — the only one available at the time of writing — explicitly acknowledges that **base two-qubit gate fidelity will likely not significantly improve over the coming years**. Progress must come from quantum error correction, not from better physical qubits. [EST] And quantum error correction carries a factor of **100,000 uncertainty** in eventual performance, according to the same roadmap. [EST]

Neutral Atoms — The Scaling Surprise

Neutral-atom quantum computers use uncharged atoms trapped in optical tweezers — focused laser beams that act as microscopic “tractor beams.” Qubits are encoded in two atomic states, and gates are performed using the Rydberg blockade effect, where exciting one atom to a highly excited state prevents its neighbors from being similarly excited. [EST]

Why they are surging: Neutral atoms have demonstrated the fastest scaling of any platform. Optical tweezer arrays can trap thousands of atoms in reconfigurable geometries. Coherence times are excellent — seconds, rivaling trapped ions. And because atoms are identical by nature, there are no fabrication variations between qubits. [EST]

In 2026, IEEE Spectrum called neutral atoms “2026’s Big Leap” — the platform best positioned to scale physical qubit counts rapidly without the fabrication bottlenecks of superconducting circuits or the shuttling constraints of trapped ions. [PROP]

Why they are not yet ahead: Gate fidelities lag behind trapped ions. The Rydberg blockade mechanism is less mature. Laser systems become extremely complex at scale — controlling thousands of individually addressed atoms requires thousands of

laser beams with nanometer precision. And like every platform, neutral atoms ultimately face the same thermodynamic wall that active error correction entails. [EST]

Photonic Quantum Computing — Room Temperature, Photon Loss

Photonic quantum computers encode information in individual particles of light — photons. Because photons interact weakly with their environment, photonic systems can operate at room temperature, eliminating the cryogenic cooling problem entirely. [EST]

Why they are appealing: Room-temperature operation. Natural compatibility with fiber-optic networks — photons are the native currency of communication. PsiQuantum has raised over \$665 million in private funding — plus a A\$940 million Australian government package announced in 2024 — to develop silicon photonic quantum computers using existing semiconductor fabrication infrastructure. [EST]

Why they are not yet proven: Photon loss is a fundamental challenge. Photons can be absorbed, scattered, or simply fail to arrive at detectors. Two-qubit gates require photons to interact, and photons do not naturally interact — they pass through each other. Creating effective photon-photon interactions requires either massive multiplexing overhead (generating many photons and post-selecting the ones that happened to interact) or strong optical nonlinearities that are difficult to achieve. No photonic platform has demonstrated a fault-tolerant logical qubit. [EST]

Silicon Spin Qubits — The Semiconductor Play

Silicon spin qubits encode information in the spin of a single electron confined in a quantum dot — essentially a nanoscale transistor. They leverage the same fabrication techniques that produce classical computer chips, using isotopically purified silicon to eliminate nuclear spin noise. [EST]

Why they are promising: Silicon spin qubits are the most “classical-friendly” platform. They are nanoscale — far smaller than superconducting qubits — enabling high density. They are manufactured using existing semiconductor fabrication infrastructure. Companies like Intel, Diraq, and Silicon Quantum Computing are pursuing this approach. [EST]

Why they are behind: Silicon spin qubits are 5-10 years behind superconducting qubits in maturity. Gate fidelities are improving rapidly but are not yet competitive with trapped ions or superconducting qubits. Charge noise and isotopic impurities remain challenges. And like superconducting qubits, they require millikelvin temperatures. [EST]

Topological Qubits — The High-Risk Bet

Topological quantum computing encodes information not in the state of a single particle but in the collective behavior of many particles arranged in a specific pattern — a topological phase of matter. The information is stored non-locally, making it intrinsically protected from local noise. [EST] The specific proposal pursued by Microsoft uses Majorana zero modes — exotic quasiparticles that are their own antiparticles — as the computational building blocks.

Why they are potentially revolutionary: If topological qubits work as theorized, they would reduce the physical-to-logical qubit overhead from the current $\sim 1,000:1$ to perhaps $10:1$. This would change the entire scaling equation. Error correction would become a modest refinement rather than an all-consuming resource drain. [PROP]

Why they are disputed: In February 2025, Microsoft announced Majorana 1 — what they called “the world’s first quantum processor powered by topological qubits.” [EST] The announcement claimed the creation of a new state of matter — a topological superconductor — and the measurement of Majorana zero modes at the ends of nanowires. But leading physicists have expressed skepticism. Scientific American reported that some physicists believe “the approach of building a quantum computer based on topological Majorana qubits as it is pursued by Microsoft is not going to work.” [EST]

The history of Majorana claims in condensed matter physics includes retractions. A 2018 Microsoft-affiliated paper in *Nature* claiming Majorana evidence was retracted in 2021. [EST] The stakes are enormous — Microsoft has invested over two decades in this approach — but whether it will work remains genuinely unknown.

The Bottom Line: Platform Diversity Is Rational

No single platform has won. None appears poised to win in the next five years. The practical recommendation for anyone building a quantum strategy in 2026 is simple: **access multiple platforms via cloud services.** Do not bet on a single modality. IBM Quantum, IonQ, QuEra, and others all offer cloud access. The cost of platform diversity — learning multiple software development kits, porting algorithms between architectures — is far lower than the cost of betting on the wrong hardware. [PROP]

But platform choice is ultimately secondary. All six platforms face the same deeper problem — the one we turn to in the next chapter.

Chapter 3: The Thermodynamic Wall — Why Error Correction Is the Real Bottleneck

Why Qubits Need Error Correction

Classical computers do not need error correction for logic operations. A transistor switches between 0 and 1 with an error rate of roughly 10^{-18} to 10^{-27} — so low that you can run a computer for years without a single bit flip from random noise. [EST] Quantum computers are fundamentally different. Qubits are analog systems — they exist in superpositions, and those superpositions are fragile. A stray photon, a thermal phonon, a fluctuating electric field — any interaction with the environment causes decoherence, destroying the quantum information.

The error rates of physical qubits are approximately 10^{-3} to 10^{-4} per gate — meaning that roughly one in a thousand to one in ten thousand gate operations produces an error. [EST] This is a billion billion times worse than classical transistors. To run a quantum algorithm — which might require millions or billions of gates — you need to correct these errors as they occur.

How Quantum Error Correction Works (The Short Version)

Quantum error correction (QEC) works by spreading quantum information across multiple physical qubits so that errors can be detected without measuring — and thereby destroying — the quantum state itself. This is the clever part: you cannot simply “look” at a qubit to check if it has an error, because looking at a qubit collapses its superposition. Instead, you perform **syndrome measurements** — measurements that reveal whether an error occurred without revealing the encoded information. [EST]

The dominant QEC architecture is the **surface code**. It arranges physical qubits on a two-dimensional grid and performs repeated syndrome measurements — patterns of neighboring qubit measurements — to detect errors. When an error is detected, a classical computer calculates the most likely correction and applies it. [EST]

The surface code works. Google demonstrated it below threshold in 2024-2025 — proving that increasing the code size (the “distance”) suppresses the logical error rate. [EST] This was a genuine milestone, pursued for over twenty years.

But “below threshold” does not mean “usefully low.” It means the logical error rate is lower than the physical error rate. To achieve error rates low enough for practical computation — say, Shor’s algorithm factoring a cryptographically relevant number — you need error rates of roughly 10^{-15} per logical gate. [PROP] The surface code achieves this by increasing the code distance — which means adding more physical qubits per logical qubit. Current estimates place the overhead at approximately **1,000 physical qubits per logical qubit**. [PROP] For a computation requiring 10,000 logical qubits, you would need 10 million physical qubits — plus the syndrome measurement infrastructure to monitor them all.

The Cooling Gap

And here is the problem. Every syndrome measurement costs energy. Every physical qubit dissipates heat. Every control line adds thermal load. All of this happens inside a dilution refrigerator operating at millikelvin temperatures.

A commercial dilution refrigerator provides approximately **50 microwatts (50 μ W) of cooling at the mixing chamber** — where the qubits live. The pulse-tube cryocooler at the 4 kelvin stage provides approximately **1 watt** — a factor of **20,000× more**. [EST]

The gap exists because moving heat from millikelvin temperatures to room temperature requires enormous energy input, constrained by the Carnot efficiency of heat pumps. The efficiency scales with the temperature ratio: to move 1 watt of heat from 20 mK to 300 K requires at minimum approximately 15 kilowatts of room-temperature power, purely from thermodynamics. With realistic efficiencies (~30% of Carnot), it climbs to approximately 50 kilowatts per watt of millikelvin cooling. [EST]

Now do the arithmetic. Each surface code logical qubit, at scale, would dissipate roughly 24 watts — not at the millikelvin stage directly, but in the control and measurement infrastructure that feeds into it. [PROP] Even a fraction of that reaching the mixing chamber would overwhelm the 50-microwatt budget. At 10,000 logical

qubits — the scale needed for Shor’s algorithm — the room-temperature wall power would approach 240 kilowatts. At 1 million logical qubits, it exceeds 24 megawatts. [EST] for the cooling physics; [PROP] for the extrapolation to quantum scale.

This is not an engineering problem. It is a **thermodynamic** problem. The 20,000× gap between what the 4 K stage can provide and what the millikelvin stage can deliver is a consequence of the Carnot limit — the fundamental bound on heat pump efficiency. No amount of engineering optimization can bridge a Carnot-limited gap of this magnitude. [EST]

A 2026 successor analysis treats the thermodynamic and informational bottlenecks jointly (10.5281/zenodo.17955898): an integrated thermal-equilibrium model of a machine designed to run Shor’s algorithm finds that even under optimistic technological assumptions, the required qubit count, cooling power, and classical decoding speed remain orders of magnitude beyond plausible near-term capabilities. The barrier is not a single engineering challenge but a multi-physics resource gap. [PROP]

The Surface Code Plateau

This is why the surface code is not a path to practical quantum computing at commercially useful scales. Google’s 2023 demonstration (*Nature*, Vol. 614, p. 676) showed that increasing the surface code distance from 3 to 5 reduced the logical error rate from approximately 3% to approximately 2.9% — a marginal improvement. [EST] Extrapolating this trend linearly to distance 11 projects error rates of approximately 2.6% — still far from useful. [PROP] (A linear extrapolation is the conservative plateau model; Google’s measured per-cycle suppression factor, **Math input error**, would project a lower rate under exponential scaling — but the qubit, measurement, and cooling costs grow with distance either way.)

The 2024-2025 demonstrations of “below threshold” QEC (*Nature*, Vol. 638, p. 920) were genuine breakthroughs. But the threshold is a mathematical line — below it, error rates decrease with code distance; above it, they increase. Being below threshold does not mean the error rate is low enough. It means you are on the right side of the line. The distance from “below threshold” to “usefully low” is orders of magnitude, and closing it requires increasing code distance, which requires more physical qubits, which requires more syndrome measurements, which requires more cooling — which hits the thermodynamic wall. [EST]

What This Means

The standard quantum computing story — the one told in most textbooks — goes like this: we are in the NISQ era; qubits are improving; error correction is being developed; and in five to ten years, we will have fault-tolerant quantum computers that change the world.

The physics tells a different story. The NISQ era is not a temporary phase on the way to fault tolerance. It is the operating regime of ALL current architectures, and the path out of it — active quantum error correction — faces a thermodynamic wall that is a hard physical limit, not an engineering challenge. [PROP]

The next part of this guide will teach you the standard quantum computing formalism — qubits, entanglement, gates, algorithms — because you need to understand it. But you will understand it with the awareness that the platform it describes may never scale to practical utility. And Part IV will introduce an alternative: a different geometry for computing, one where error correction is passive — built into the hardware structure itself — rather than active, measurement-intensive, and thermodynamically unsustainable.

For now, the key insight is this: **the bottleneck in quantum computing is not qubit count. It is not gate fidelity. It is not algorithm design. It is thermodynamics.** And thermodynamics does not negotiate.

Chapter Summary — Part I

Chapter	Key Insight	Confidence
1. Why Don't We Have Quantum Computers?	The NISQ era has lasted longer than expected; no commercially useful quantum computation has been demonstrated	[EST]
2. The Six Platforms	No single platform has won; all face fundamental scaling challenges; platform diversity is rational	[EST]
3. The Thermodynamic Wall	Active QEC faces a 20,000× cooling gap that is Carnot-limited — a hard physical bound, not an engineering problem	[EST] for the cooling physics; [PROP] for extrapolation to QEC at scale

References — Part I

- Preskill, J. (2018). “Quantum Computing in the NISQ era and beyond.” *Quantum*, 2, 79. [EST]
- Hempel, C. (2026). “Trapped-Ion Quantum Computers.” In Jang-Jaccard et al. (eds.), *Quantum Technologies*. Springer. DOI: 10.1007/978-3-031-90727-2_2. [EST]
- Acharya, R. et al. (2025). “Quantum error correction below the surface code threshold.” *Nature*, 638, 920-926. [EST]
- Google Quantum AI. (2023). “Suppressing quantum errors by scaling a surface code logical qubit.” *Nature*, 614, 676-681. [EST]
- Quni-Gudzinis, R. B. (2025). *Thermodynamic Constraints and Architectural Inversions in Scalable Quantum Information Systems*. Zenodo. DOI: 10.5281/zenodo.17938113. [PROP]
- Quni-Gudzinis, R. B. (2025). *The Lifecycle of a Fault-Tolerant Quantum Computer*. Archive: 2025/12. [PROP]
- Quni-Gudzinis, R. B. (2025). *Thermodynamic and Informational Bottlenecks of Scalable Fault-Tolerant Quantum Computation*. Zenodo. DOI: 10.5281/zenodo.17955898. [PROP]

End of Part I. Part II (Chapters 4-7) will cover the standard quantum computing curriculum: qubits, superposition, entanglement, quantum algorithms, and quantum error correction — taught clearly, with the awareness developed in Part I.

Part II: The Standard Curriculum

Chapter 4: Qubits, Superposition, and the Bloch Sphere

The Classical Bit vs. the Quantum Bit

Every computer you have ever used is built from bits — two-state systems that are either 0 or 1. A bit is physically implemented as a voltage level (high or low), a magnetic domain (up or down), or a charge on a capacitor (charged or discharged). The defining property of a classical bit is **determinism**: at any moment, a bit has a definite value. [EST]

A quantum bit — a **qubit** — is different. A qubit can exist in a **superposition** of 0 and 1: a state that is partly 0 and partly 1. This is not a metaphor. It is a physical fact, verified in thousands of experiments. [EST]

The most common mathematical representation uses two basis states:

$$|0\rangle = \begin{pmatrix} 1 \\ 0 \end{pmatrix}, \quad |1\rangle = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$$

The notation $|\cdot\rangle$ is called a **ket** — Paul Dirac’s notation for quantum states. The most general state of a single qubit is:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle = \begin{pmatrix} \alpha \\ \beta \end{pmatrix}$$

where α and β are complex numbers satisfying $|\alpha|^2 + |\beta|^2 = 1$. This normalization condition ensures that the total probability of finding the qubit in some state is 1 — a physical requirement. [EST]

What Superposition Means (and Doesn’t Mean)

Superposition is often explained with phrases like “the qubit is both 0 and 1 at the same time.” This is catchy but misleading. A more precise statement: **before measurement, the qubit does not have a definite value. The act of measurement forces it into either $|0\rangle$ with probability $|\alpha|^2$ or $|1\rangle$ with probability $|\beta|^2$.** [EST]

This is the **Born rule**, named after Max Born who proposed it in 1926. It is one of the foundational postulates of quantum mechanics, and it has been experimentally confirmed to extraordinary precision. [EST]

Superposition does NOT mean: - The qubit is secretly 0 or 1 and we just don't know which. (That would be a classical probability distribution, not a quantum superposition.) - The qubit is literally in two places at once. (It is in one physical location — it is the *information* that is distributed.) - You can “look” at a superposition without disturbing it. (Measurement is destructive: it collapses the superposition into a definite outcome.)

What superposition DOES mean: the qubit carries information in a way that has no classical analog. Two complex numbers (α and β), constrained by one real condition ($|\alpha|^2 + |\beta|^2 = 1$), specify a point on a three-dimensional sphere. That is a lot of information in a single physical system — and it is this information density that makes quantum computing potentially powerful. [EST]

The Bloch Sphere — A Geometric Picture

The state of a single qubit can be visualized on the **Bloch sphere**, a unit sphere in three dimensions. Any pure state $|\psi\rangle$ corresponds to a point on the surface:

$$|\psi\rangle = \cos\left(\frac{\theta}{2}\right)|0\rangle + e^{i\phi}\sin\left(\frac{\theta}{2}\right)|1\rangle$$

where $\theta \in [0, \pi]$ is the polar angle (latitude) and $\phi \in [0, 2\pi)$ is the azimuthal angle (longitude). [EST]

Key points on the Bloch sphere: - North pole ($\theta = 0$): $|0\rangle$ - South pole ($\theta = \pi$): $|1\rangle$ - Equator ($\theta = \pi/2$, any ϕ): equal superpositions, e.g., $(|0\rangle + |1\rangle)/\sqrt{2}$

The Bloch sphere makes single-qubit operations intuitive: a quantum gate is simply a rotation of the Bloch sphere. This geometric picture is one of the most elegant and useful tools in quantum computing — but it only works for a single qubit. For multiple qubits, the geometry becomes vastly more complex. [EST]

Why This Matters (and What It Doesn't Guarantee)

The information density of a qubit is genuinely remarkable. A single qubit's state lives on a continuous sphere — in principle, it stores an infinite amount of information in the real numbers θ and ϕ . In practice, you cannot extract that information. A measurement gives you one bit: 0 or 1. The power of quantum computing comes not from storing information in a single qubit but from **entanglement** — correlations between multiple qubits that have no classical analog. That is the subject of the next chapter. [EST]

Chapter 5: Entanglement and Multi-Qubit Systems

Combining Qubits — The Tensor Product

When you have two classical bits, the possible states are 00, 01, 10, and 11 — four possibilities. With two qubits, the state space is the **tensor product** of the individual spaces, spanned by four basis states: $|00\rangle$, $|01\rangle$, $|10\rangle$, $|11\rangle$. A general two-qubit state is:

$$|\psi\rangle = \alpha_{00}|00\rangle + \alpha_{01}|01\rangle + \alpha_{10}|10\rangle + \alpha_{11}|11\rangle$$

with $\sum |\alpha_{ij}|^2 = 1$. [EST]

For n qubits, the state requires 2^n complex numbers to describe — exponential growth. This is Feynman’s original insight: simulating n quantum particles on a classical computer requires resources that grow exponentially with n , while a quantum computer of n qubits does not. [EST]

Entanglement — Correlations Beyond Classical Physics

Some two-qubit states can be written as a product of single-qubit states: $|\psi\rangle = |a\rangle \otimes |b\rangle$. These are called **separable** or **product** states. But most two-qubit states cannot be factored this way. These are **entangled** states. [EST]

The canonical example is the **Bell state** (named after John Stewart Bell):

$$|\Phi^+\rangle = \frac{|00\rangle + |11\rangle}{\sqrt{2}}$$

If you measure the first qubit of $|\Phi^+\rangle$ and get 0, the second qubit is instantaneously determined to be 0 — regardless of how far apart the qubits are. If you get 1, the second qubit is 1. The outcomes are perfectly correlated, even though neither qubit had a definite value before measurement. [EST]

This is the phenomenon Einstein derisively called “spooky action at a distance” — he intended the phrase as a criticism, arguing that quantum mechanics must be incomplete if it permitted such correlations. Subsequent experiments, most conclusively the 2015 loophole-free Bell tests, have confirmed that entanglement is real and that Einstein’s alternative (local hidden variables) is ruled out. The “spookiness” is genuine, but it does not violate the no-signaling theorem: entanglement cannot be used for faster-than-light communication. It has been experimentally verified in thousands of experiments since the 1970s, most conclusively in the 2015 “loophole-free Bell tests” that closed every remaining experimental loophole. [EST]

What entanglement does NOT mean: - It does not allow faster-than-light communication. (You cannot control which outcome you get, so you cannot send a message.) - It does not mean the particles are “connected by a force.” (The correlation is non-local but does not transmit energy or information.)

What entanglement DOES mean: Quantum systems can exhibit correlations that are stronger than any possible classical correlation. These correlations are the fuel of quantum computation — they are what enable quantum algorithms to do things that classical algorithms cannot. [EST]

Quantum Gates — Operations on Qubits

A quantum gate is a linear transformation that preserves the normalization of the state — mathematically, a **unitary matrix** U satisfying $U^\dagger U = I$. [EST]

The most important single-qubit gates:

Gate	Symbol	Matrix	Effect
Pauli-X	X	$\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$	Bit flip: $ 0\rangle \leftrightarrow 1\rangle$
Pauli-Z	Z	$\begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}$	Phase flip: $ 1\rangle \rightarrow - 1\rangle$
Hadamard	H	$\frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$	Creates superposition: $ 0\rangle \rightarrow (0\rangle + 1\rangle)/\sqrt{2}$

The most important two-qubit gate is the **CNOT** (controlled-NOT):

$$\text{CNOT} = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

CNOT flips the target qubit if and only if the control qubit is $|1\rangle$. It is the gate that creates entanglement from separable states. [EST]

Any multi-qubit unitary operation can be decomposed into a sequence of single-qubit gates and CNOT gates. This set — $\{H, \text{CNOT}, T\}$ where T is a specific phase gate — is called a **universal gate set**. It can approximate any quantum computation to arbitrary precision. [EST]

The Quantum Circuit Model

A quantum circuit is a sequence of quantum gates applied to qubits, followed by measurement. It is the quantum analog of a classical logic circuit. Circuits are read left to right: qubits start in the $|0\rangle$ state, gates are applied in sequence, and the final state is measured. [EST]

The circuit model is the most common framework for describing quantum algorithms, but it is not the only one. **Measurement-based quantum computing** (MBQC) starts with a large entangled state and performs computation through a sequence of measurements. **Adiabatic quantum computing** encodes the solution in the ground state of a slowly varying Hamiltonian. These alternative models are computationally equivalent to the circuit model — any algorithm expressed in one can be translated to the others — but they suggest different hardware implementations. [EST]

Chapter 6: Quantum Algorithms — The Promise

The Idea of Quantum Speedup

The reason anyone cares about quantum computing is the promise of **quantum speedup** — solving certain problems faster than any classical computer can. The speedup is not universal. Quantum computers are not “faster computers” for all

problems. They provide asymptotic advantages for specific problem classes where quantum phenomena — superposition, entanglement, interference — can be exploited to avoid the exponential scaling that classical algorithms face. [EST]

We will examine three algorithms that illustrate the spectrum of quantum speedup: Deutsch-Jozsa (exponential advantage for a contrived problem), Grover's search (quadratic advantage), and Shor's factoring (exponential advantage for a practically important problem). [EST]

Deutsch-Jozsa — The Simplest Quantum Speedup

The Deutsch-Jozsa algorithm solves a deliberately artificial problem: determine whether an unknown function $f : \{0, 1\}^n \rightarrow \{0, 1\}$ is **constant** (same output for all inputs) or **balanced** (outputs 0 for exactly half the inputs and 1 for the other half). A classical algorithm might need up to $2^{n-1} + 1$ queries in the worst case. The Deutsch-Jozsa algorithm solves it with a single query — an exponential speedup. [EST]

This is genuine but limited. The problem is contrived. No one needs to solve Deutsch-Jozsa in practice. Its value is pedagogical: it demonstrates that quantum speedup is possible, using the simplest nontrivial case. [EST]

Grover's Search — Quadratic Speedup

Grover's algorithm searches an unstructured database of N items in $O(\sqrt{N})$ steps, compared to $O(N)$ for classical search. For $N = 1,000,000$, Grover's requires approximately 1,000 steps rather than 500,000 (average) — a speedup of $500\times$. [EST]

The algorithm works by **amplitude amplification**: starting from a uniform superposition over all items, it repeatedly applies an operation that increases the amplitude of the marked solution while decreasing all others. After approximately $\pi\sqrt{N}/4$ iterations, the solution's amplitude is close to 1, and measurement yields the answer with high probability. [EST]

Grover's algorithm provides a **quadratic** speedup — useful but not transformative. A problem that takes a classical computer one year would take a quantum computer running Grover's about two weeks. This is significant but not the kind of exponential advantage that would revolutionize computation. [EST]

Grover's also has implications for cryptography: it can search for a symmetric encryption key (e.g., a 128-bit AES key) in approximately 2^{64} steps rather than 2^{128} . This is why the cryptographic community recommends doubling symmetric key lengths to prepare for quantum computers: AES-128 becomes roughly as secure against Grover's as a classical 64-bit key would be. [EST]

Shor's Algorithm — The Exponential Revolution (in Theory)

Shor's algorithm factors an n -bit integer in $O(n^3)$ quantum operations, compared to the best known classical algorithm (the general number field sieve) which runs in roughly $\exp(O(n^{1/3}))$ time. For $n = 2048$ (the size of an RSA key), Shor's would require approximately 10^9 to 10^{10} quantum gate operations — a large but finite number. [EST]

Shor's algorithm works by reducing factoring to **period finding** and solving the period-finding problem using the **quantum Fourier transform** (QFT). The QFT is the quantum analog of the classical fast Fourier transform and is the engine behind Shor's exponential speedup. [EST]

The catch — which most guides don't emphasize enough: Shor's algorithm requires **fault-tolerant logical qubits** — qubits with error rates low enough that the computation can run to completion without accumulating fatal errors. At current physical error rates ($\sim 10^{-3}$) and surface code overhead ($\sim 1,000$ physical qubits per logical qubit), factoring a 2048-bit RSA key would require roughly **20 million physical qubits** — far beyond any existing or near-term device. [PROP]

This is the reality check. Shor's algorithm is mathematically beautiful and algorithmically revolutionary. Its physical realization — at cryptographically relevant scale — requires solving the thermodynamic wall problem described in Chapter 3. Whether that problem is solvable within the standard (active QEC) paradigm is the central open question. [PROP]

The Gate Decomposition Problem

Every quantum algorithm expressed in high-level gates (like those above) must be decomposed into the native gate set of a specific hardware platform. For example, the Clifford+T gate set — $\{H, S, \text{CNOT}, T\}$ where T is a $\pi/4$ phase rotation — is universal. The T gate is the expensive one: it requires magic state distillation, a procedure that consumes many physical qubits per T gate. [EST]

The Kliuchnikov-Maslov-Mosca algorithm (2013) provides an optimal decomposition of single-qubit unitaries into Clifford+T gates with an exact optimality guarantee on the number of T gates. [EST] But even optimal decomposition does not change the fundamental overhead: every T gate costs resources, and fault-tolerant Shor's algorithm requires millions of them. [PROP]

Chapter 7: Quantum Error Correction — How It's Supposed to Work

Why Classical Error Correction Fails

Classical error correction works by **redundancy**: store each bit multiple times and use majority voting. If you store 0 as 000 and one bit flips to 001, you can correct it back to 000. This works because you can **measure** the bits without disturbing them — reading a classical bit does not change its value. [EST]

Quantum error correction faces two fundamental obstacles that classical ECC does not:

1. **The no-cloning theorem:** You cannot copy an arbitrary quantum state. There is no quantum analog of “store the bit three times.” [EST]
2. **Measurement destroys superposition:** You cannot simply “look” at a qubit to see if it has an error, because looking collapses the superposition. [EST]

Stabilizer Codes — Detecting Errors Without Measuring the State

The solution is remarkably elegant: instead of copying the state, **spread the quantum information across multiple physical qubits** and design measurements that reveal whether an error occurred without revealing the encoded information. [EST]

The mathematical framework is the **stabilizer formalism**. For an $[[n, k, d]]$ code, n physical qubits encode k logical qubits with code distance d (the minimum number of physical errors that cause a logical error). The code is defined by $n - k$ stabilizer generators — multi-qubit operators whose measurement reveals error syndromes. [EST]

The simplest nontrivial code is the **3-qubit bit-flip code** ($[[3, 1, 1]]$):

$$|0_L\rangle = |000\rangle, \quad |1_L\rangle = |111\rangle$$

Stabilizer generators: $Z_1 Z_2$ and $Z_2 Z_3$ (measuring these detects X errors on qubits 1 or 2, and 2 or 3 respectively). This code corrects a single bit-flip error but cannot correct phase-flip errors. [EST]

To correct both X and Z errors, you need the **9-qubit Shor code** ($[[9, 1, 3]]$), which concatenates the bit-flip and phase-flip codes. This was the first quantum error-correcting code ever discovered (Shor, 1995). [EST]

The Surface Code — The Workhorse of QEC

The surface code is the dominant QEC architecture. It arranges physical qubits on a two-dimensional grid, with data qubits (storing information) interspersed with syndrome qubits (measuring errors). Two types of syndrome measurements — X -stabilizers (detecting Z errors) and Z -stabilizers (detecting X errors) — are performed repeatedly. [EST]

The surface code is appealing because: - It requires only nearest-neighbor interactions — no long-range connectivity. - It has a relatively high threshold — the physical error rate below which increasing code distance reduces logical error. - It is well-understood theoretically and has been demonstrated experimentally. [EST]

The surface code's costs: - **Encoding rate vanishes:** As code distance increases, $k/n \rightarrow 0$. The fraction of qubits doing useful computation goes to zero. - **Decoding is expensive:** Minimum-Weight Perfect Matching (MWPM) decoding requires $O(N^3)$ classical computation. - **Syndrome measurements are continuous:** Every gate cycle requires syndrome readout, generating heat. - **Physical overhead is enormous:** Approximately 1,000 physical qubits per logical qubit for practically useful error rates. [EST]

Magic State Distillation — The Hidden Cost

The Clifford gates (H, S, CNOT) can be implemented fault-tolerantly on the surface code with relatively modest overhead. But the T gate — required for universal quantum computation — cannot. It requires **magic state distillation**: a procedure that consumes many noisy copies of a special state (the “magic state”) and produces one high-fidelity copy. [EST]

Magic state distillation is the dominant resource cost in fault-tolerant quantum computing. For Shor’s algorithm factoring a 2048-bit number, estimates place the T-gate count in the billions, with each T gate requiring hundreds to thousands of physical qubits for distillation. [PROP]

This is the arithmetic that underlies the thermodynamic wall argument from Chapter 3. Every gate — especially every T gate — consumes physical resources. Every physical resource generates heat. The heat must be removed by cryogenic systems operating at millikelvin temperatures. And the Carnot limit imposes a fixed ratio between heat removed at millikelvin temperatures and power consumed at room temperature. [EST]

Where We Are Now

The surface code has been demonstrated below threshold — a genuine milestone. Logical qubits with suppressed error rates have been created. Real-time error correction feedback loops have been closed (Quantinuum, 2024). [EST]

But “below threshold” is not the same as “usefully low.” The gap between demonstrated error rates ($\sim 10^{-2}$ per logical gate) and practically useful error rates ($\sim 10^{-15}$) is roughly thirteen orders of magnitude. Closing this gap through increased code distance means adding physical qubits, which means adding syndrome measurements, which means adding heat — which brings us back to the thermodynamic wall. [PROP]

Chapter Summary — Part II

Chapter	Key Insight	Confidence
4. Qubits and Superposition	A qubit is a two-state quantum system; superposition is real but measurement is destructive; the Bloch sphere provides geometric intuition	[EST]
5. Entanglement and Gates	Entanglement enables correlations beyond classical physics; quantum gates are unitary matrices; CNOT + single-qubit gates = universal	[EST]
6. Quantum Algorithms	Speeds vary: Deutsch-Jozsa (exponential, contrived), Grover’s (quadratic, useful), Shor’s (exponential, revolutionary — but requires fault tolerance at unbuilt scale)	[EST] for algorithms; [PROP] for resource estimates
7. Quantum Error Correction	Surface codes work below threshold; encoding rate vanishes; T-gate distillation	[EST] for QEC theory; [PROP] for scaling extrapolation

Chapter	Key Insight	Confidence
	dominates resource costs; thermodynamic wall constrains scaling	

References — Part II

1. Nielsen, M. A., & Chuang, I. L. (2010). *Quantum Computation and Quantum Information*. Cambridge University Press. [EST] — The standard textbook.
2. Preskill, J. (2018). “Quantum Computing in the NISQ era and beyond.” *Quantum*, 2, 79. [EST]
3. Shor, P. W. (1994). “Algorithms for quantum computation: discrete logarithms and factoring.” *FOCS 1994*. [EST]
4. Grover, L. K. (1996). “A fast quantum mechanical algorithm for database search.” *STOC 1996*. [EST]
5. Fowler, A. G., Mariantoni, M., Martinis, J. M., & Cleland, A. N. (2012). “Surface codes: Towards practical large-scale quantum computation.” *Physical Review A*, 86(3), 032324. [EST]
6. Kliuchnikov, V., Maslov, D., & Mosca, M. (2013). “Fast and Efficient Exact Synthesis of Single-Qubit Unitaries Generated by Clifford and T Gates.” *Quantum Information and Computation*, 13(7&8), 607-630. [EST]
7. Acharya, R. et al. (2025). “Quantum error correction below the surface code threshold.” *Nature*, 638, 920-926. [EST]
8. Quni-Gudzinis, R. B. (2025). *Thermodynamic Constraints and Architectural Inversions in Scalable Quantum Information Systems*. Zenodo. DOI: 10.5281/zenodo.17938113. [PROP]

End of Part II. Part III (Chapters 8-10) will cover The Reality Check: the surface code plateau, the cryogenic arithmetic, and what quantum computing actually delivers today.

Part III: The Reality Check

Chapter 8: The Surface Code Plateau

What “Below Threshold” Actually Means

In 2023, Google Quantum AI published a landmark experiment demonstrating that increasing the surface code distance from $d = 3$ to $d = 5$ reduced the logical error rate — proving that error correction works. [EST] The result was celebrated as a major milestone, which it was. But the numbers tell a more sober story.

The logical error rate at $d = 3$ was approximately 3%. At $d = 5$, it was approximately 2.9%. [EST] The improvement was real but marginal — a reduction of roughly 0.1 percentage points for a doubling of code distance. Extrapolating this trend linearly to $d = 11$ — the distance often cited as the threshold for “usefully low” error rates — projects a logical error rate of approximately 2.6%. [PROP] (A

linear extrapolation is the conservative plateau model; Google’s measured per-cycle suppression factor, **Math input error**, would project a lower rate under exponential scaling.)

Two-point-six percent. That is the error rate per logical gate that the best-understood error correction architecture projects at practically achievable code distances under the conservative linear extrapolation. For context, a computation involving a million logical gates — a modest quantum algorithm — would experience roughly 26,000 errors. The computation would fail. [EST]

The Threshold Is a Line, Not a Destination

In 2024-2025, Google and others demonstrated quantum error correction **below the surface code threshold**. [EST] This was the milestone the field had pursued for over twenty years. The threshold is a specific physical error rate — approximately 1% for depolarizing noise under standard decoding — below which increasing code distance suppresses logical errors rather than amplifying them. Being below threshold means the error correction is working in the right direction: error rates decrease with code size. Being above threshold means they increase — error correction makes things worse. [EST]

Below threshold is genuine progress. But it is a line you cross, not a destination you arrive at. “Below threshold” means the logical error rate at $d = 5$ is lower than at $d = 3$. It does not mean the logical error rate is low enough to run Shor’s algorithm. It does not mean the error rate is low enough to simulate a molecule. It means you are on the right side of a mathematical inequality — and the distance from that line to useful computation remains vast. [EST]

The gap between demonstrated error rates ($\sim 10^{-2}$) and the error rates needed for practical computation ($\sim 10^{-15}$) is roughly thirteen orders of magnitude — a factor of 10,000,000,000,000. Closing this gap by increasing code distance is theoretically possible. But each increase in code distance requires more physical qubits, more syndrome measurements, more classical decoding, and more energy dissipation. [PROP]

The Surface Code vs. The Tree Code — A Preview

We will explore the alternative in detail in Part IV. For now, a preview of why alternative code geometries matter. The surface code and the ultrametric tree code differ not just in their error thresholds but in their fundamental resource scaling:

Property	Surface Code	Tree Code (BTQP)
Bit-flip threshold	$\sim 10.9\%$	50.0% (4.6×) [PROP]
Depolarizing threshold	$\sim 1.0\%$	75.0% (75×) [PROP]
Physical qubits for $d = 11$ equivalent	241	~ 40 [PROP]
Decoding complexity	$O(N^3)$ (global)	$O(s \log s)$ (local, parallel) [PROP]
Encoding rate	$k/n \rightarrow 0$ (vanishes)	$1 - 1/p$ (constant; $\sim 1/2$) [PROP]

Property	Surface Code	Tree Code (BTQP)
Syndrome measurements	Active (120/cycle)	Passive (0) [PROP]

These numbers come from the Bruhat-Tits Quantum Processor paper [PROP] — computationally validated, not yet experimentally demonstrated. But they illustrate why “the surface code is working, just give it time” may not be the right conclusion from the data. [PROP]

Chapter 9: The Cryogenic Arithmetic

How a Dilution Refrigerator Works

A dilution refrigerator is the device that cools superconducting and spin qubits to their operating temperature of approximately 10-20 millikelvin — colder than interstellar space. It works by exploiting the unique properties of helium-3 and helium-4 mixtures at ultra-low temperatures. When a mixture of ^3He and ^4He is cooled below approximately 0.87 K, it separates into two phases: a ^3He -rich phase and a ^3He -dilute phase. Pumping ^3He from the dilute phase into the rich phase absorbs heat — the same principle as evaporative cooling, but exploiting quantum statistics rather than latent heat. [EST]

A commercial dilution refrigerator is a marvel of cryogenic engineering. It weighs several tons. It costs \$500,000 to \$1,000,000. It runs on a closed helium-3/helium-4 mixture — helium-3 is a scarce and expensive resource — with continuous operation of compressors, pumps, and control systems. [EST]

And it delivers approximately **50 microwatts of cooling at the mixing chamber**. [EST]

Fifty microwatts. To put that in perspective: a typical LED indicator light consumes about 20,000 microwatts. The qubits in a dilution refrigerator must dissipate less heat — combined, across all control lines, all measurements, all wiring — than a single LED. [EST]

The Carnot Limit

The Carnot efficiency is the maximum possible efficiency of any heat engine operating between two temperatures. For a refrigerator — a heat engine run in reverse — the Carnot coefficient of performance is:

$$\text{COP}_{\text{Carnot}} = \frac{T_{\text{cold}}}{T_{\text{hot}} - T_{\text{cold}}}$$

For $T_{\text{cold}} = 20 \text{ mK} = 0.02 \text{ K}$ and $T_{\text{hot}} = 300 \text{ K}$:

$$\text{COP}_{\text{Carnot}} = \frac{0.02}{300 - 0.02} \approx 6.7 \times 10^{-5}$$

This means that to remove 1 watt of heat from the mixing chamber at 20 mK, you must supply at minimum approximately **15,000 watts of room-temperature power** — purely from the laws of thermodynamics. [EST]

Real refrigerators operate at a fraction of Carnot efficiency — typically 20-30%. At 30% of Carnot, removing 1 watt from 20 mK requires approximately 50,000 watts at the wall. Every milliwatt of heat dissipated inside the mixing chamber costs 50 watts of room-temperature power to remove. [EST]

The Twenty-Thousand-Fold Gap

The dilution refrigerator’s mixing chamber provides approximately 50 microwatts of cooling. The pulse-tube cryocooler at the 4 kelvin stage — just a few centimeters away — provides approximately 1 watt. The ratio: **20,000×**. [EST]

This gap is a direct consequence of the Carnot limit. The pulse-tube stage operates at 4 K — $75\times$ colder than room temperature, requiring “only” about 75 watts of room-temperature power per watt of cooling. The dilution stage operates at 0.02 K — $15,000\times$ colder. The gap between these two cooling stages is fundamental physics, not an engineering optimization problem. [EST]

Now add quantum computing. A circuit with 10,000 logical qubits using surface code error correction — at the scale needed for cryptographically relevant factoring — would dissipate approximately 240 kilowatts at room temperature, according to estimates from the QWAV research program. [PROP] For 1 million logical qubits, the power exceeds 24 megawatts. These are not extrapolations from current devices — they follow from the known energy cost of syndrome measurement, multiplied by the number of measurements required per logical qubit per gate cycle. [PROP]

The Real-World Precedent

For scale: the Large Hadron Collider at CERN operates the world’s largest cryogenic system, using 130 tons of liquid helium to cool 36,000 tons of magnets to 1.9 K — warmer than a dilution refrigerator’s mixing chamber by a factor of roughly 100. The LHC’s cryogenic system consumes tens of megawatts of electrical power. [EST]

A surface-code quantum computer at commercially useful scale would need to operate at 20 mK — roughly one hundred times colder than the LHC — while dissipating more power at the wall. This is not an apples-to-apples comparison (the LHC cools magnets, not qubits), but it illustrates the scale of the challenge. Cryogenics at the scale needed for fault-tolerant quantum computing does not exist, and building it faces physical constraints — not just cost constraints. [PROP]

From the Wall to a Benchmark: Joules per Solution

The thermodynamic wall can be restated as a measurable engineering metric. Define **joules per solution** (JPCUB): the total wall energy consumed per correct, useful answer — including cryogenics, control electronics, classical decoding, and error-correction overhead. The question this guide has been asking qualitatively — what does a correct quantum answer cost? — becomes a number any platform can be measured against.

The metric is physics-grounded rather than arbitrary. Landauer’s principle bounds the energy per erased bit from below; the Margolus–Levitin theorem bounds the speed of dynamical evolution per unit energy; Bremermann’s limit bounds computation per unit energy. A joules-per-solution measurement inherits these ceilings: a platform cannot win by spending more energy, only by answering correctly for less. [EST]

The first system-level comparison published under this benchmark (10.5281/zenodo.21880104) compares a qudit architecture against seventeen conventional qubit platforms and, under its model assumptions, projects the qudit design below 10^{-2} joules per solution. The paper’s own caveat is the one this guide insists on: the framework has not been externally validated, and the estimates are hypotheses, not demonstrated results. [PROP]

For the reader, the benchmark converts this chapter’s argument into a falsifiable prediction. Chapter 14’s tree-code numbers (241 vs ~40 physical qubits at distance-11-equivalent protection; a $75\times$ threshold advantage; zero syndrome measurements) imply a joules-per-solution claim: a passive tree-topology processor should spend its energy budget at the 4 K stage, not the 20 mK stage. When a tree-topology device exists, that prediction can be measured directly — the same way the qudit estimate was. The wall stops being an argument and becomes a number. [PROP]

Counterarguments — What the Mainstream QC Community Would Say

A critic reading this chapter would raise several objections. Engaging with them honestly strengthens the argument — and helps the reader evaluate the evidence.

Objection 1: “LDPC codes will reduce the overhead.” Low-density parity-check (LDPC) codes and other non-surface-code QEC architectures promise better encoding rates than the surface code, potentially reducing the physical-to-logical qubit ratio below 1,000:1. [EST] IBM’s qLDPC codes demonstrated improved rates in 2024. However: (a) LDPC codes still require active syndrome measurement and decoding — they reduce but do not eliminate the thermodynamic problem. (b) Even a $10\times$ improvement in encoding rate (100:1 instead of 1,000:1) would reduce the 240 kW figure to 24 kW for 10,000 logical qubits — still a massive cryogenic engineering challenge. (c) LDPC codes have stricter connectivity requirements.

Objection 2: “Higher-temperature qubits will eliminate the cooling problem.” Some qubit platforms (photonic, certain spin qubits, topological) aim to operate at higher temperatures — 4 K, 77 K, or even room temperature. If successful, these would bypass the Carnot-limited 20 mK cooling bottleneck. However: (a) No higher-temperature qubit platform has demonstrated fault-tolerant operation. (b) Photonic qubits face photon loss challenges. (c) Topological qubits are unproven. (d) Even at 4 K, the Carnot COP is ~ 0.013 , requiring ~ 75 W of room-temperature power per watt of cooling.

Objection 3: “Special-purpose quantum computers will be useful before general-purpose ones.” This is likely true and is acknowledged in Chapter 10. The thermodynamic critique applies primarily to universal, gate-model, error-corrected quantum computers running algorithms like Shor’s.

Objection 4: “The 24 W per logical qubit estimate is too high.” The 24 W figure [PROP] is an estimate based on the energy cost of syndrome measurement, classical decoding, control electronics, and cryogenic overhead for a surface-code logical qubit at scale. The actual number could be 10× lower or 10× higher. The key question is not the precise wattage but whether ANY plausible number can stay within the Carnot-limited cooling budget. Even 1 W per logical qubit at 20 mK would cost ~50 kW at the wall, making a 10,000-logical-qubit machine a 500 MW facility. [EST] for the Carnot arithmetic; [PROP] for the extrapolation.

Objection 5: “Reversible computing could reduce heat dissipation.” In principle, quantum gates are unitary and therefore reversible. In practice, initialization, measurement, and error correction are irreversible operations that DO dissipate heat. The Landauer limit sets a lower bound of $k_B T \ln 2 \approx 1.9 \times 10^{-25}$ J per bit erasure at 20 mK — but real devices operate far above the Landauer limit.

The bottom line: The thermodynamic wall is not a proof that quantum computing is impossible. It is an argument that the standard active-QEC paradigm faces hard physical constraints that are not acknowledged in most roadmaps. The counterarguments above are valid — some combination of better codes, higher-temperature qubits, and reversible computing might push the wall back. But the uncertainty is large (the trapped-ion community’s own roadmap acknowledges a factor of 100,000 uncertainty in QEC performance [EST]), and the consequences of hitting the wall are severe. The reader should evaluate the evidence, not take either side on faith.

Counterarguments — Objections to the Ultrametric Alternative

The same standard of scrutiny applied to standard QC should be applied to the proposed alternative. A critic of the ultrametric framework would raise the following objections:

Objection 6: “The tree code has never been demonstrated on quantum hardware.” This is correct. The BTQP thresholds (50% bit-flip, 75% depolarizing, 17.3% X+Z) were verified by classical simulation — Gottesman-Knill stabilizer simulation and Monte Carlo methods — not by physical quantum devices. [GAP] The computational validation shows that the mathematical mechanism works in principle. Whether it works on real qubits with realistic noise models (correlated errors, 1/f noise, leakage) is unknown. The three falsifiable predictions (E1–E3, Chapters 16–17) are designed to test exactly this.

Objection 7: “Perfect tensors may not exist for $p > 2$.” The tree code construction for the binary tree ($p = 2$, the simplest case) uses the $[[3,1,1]]$ perfect tensor, which is well-established. For larger primes ($p = 3, 5, \dots$), the existence of perfect tensors with the required properties is not proven. [OPEN] This is an acknowledged open mathematical problem. The binary-tree case is sufficient for a proof-of-concept, but scaling to higher branching factors may require new mathematical results.

Objection 8: “Tree automorphism gates may not be universal.” The BTQP architecture proposes using tree automorphisms — permutations of the tree that preserve its structure — as logical quantum gates. Whether the set of gates generated by tree automorphisms is universal (i.e., can approximate any quantum computation) has not been fully characterized. [OPEN] If tree automorphisms are not universal, the architecture would need to be supplemented with additional gate operations, potentially increasing overhead.

Objection 9: “Building a physical tree-topology qubit array is an unsolved engineering challenge.” No laboratory has built a Bruhat-Tits tree topology with physical qubits at the leaves and majority-vote error correction at internal vertices. The five candidate platforms identified in the UQC MVP document (NV centers, neutral atoms, trapped ions, superconducting circuits, twisted superconductors) have not been experimentally tested with tree geometry. [GAP] This is a hardware engineering challenge of unknown difficulty.

Objection 10: “The thermodynamic wall argument may not apply equally to all platforms.” Photonic qubits operate at room temperature. Certain spin qubits may operate at 4 K or higher. Topological qubits, if they work, may dramatically reduce overhead. If any of these platforms succeeds at scale, the thermodynamic critique of surface-code-based architectures remains valid but less relevant. This is acknowledged in the counterarguments to the standard QC critique (Objections 1–3 above).

The bottom line on the alternative: The ultrametric framework is more speculative than the critique of standard QC. The critique of standard QC is based on well-established thermodynamics (the Carnot limit) and publicly available data (dilution refrigerator specs, Google’s error correction results). The ultrametric alternative is based on a novel mathematical framework, classically validated but not experimentally tested, with several acknowledged open problems. A reader should assign lower confidence to the positive thesis than to the negative critique — and both theses should be evaluated against future experimental evidence, particularly the three falsifiable predictions (E1–E3).

Chapter 10: What Quantum Computing Actually Delivers Today

The previous chapters have been critical of quantum computing’s failure to deliver on its transformative promises. But quantum technologies are not a monolith. Some quantum technologies deliver practical value **today**, without error correction, without dilution refrigerators, and without the thermodynamic wall. This chapter distinguishes what works now from what doesn’t.

Quantum Sensing — The Underrated Success Story

Quantum sensors exploit the same quantum properties as quantum computers — superposition, entanglement, quantum coherence — but they do not require error correction. They measure quantities (magnetic fields, electric fields, gravity, time) with precision that classical sensors cannot match. [EST]

The most mature quantum sensing technologies:

Technology	Physical System	Applications	Status
SQUIDS (Superconducting Quantum Interference Devices)	Superconducting loops with Josephson junctions	Magnetometry, brain imaging (MEG), geological survey, non-destructive testing	Deployed commercially for decades [EST]
NV centers (Nitrogen-	Single atomic defects in diamond	Nanoscale magnetometry,	Laboratory demonstrations;

Technology	Physical System	Applications	Status
Vacancy centers in diamond)		temperature sensing, biosensing	early commercial products [EST]
Atomic magnetometers	Laser-probed atomic vapors	Medical imaging, navigation, fundamental physics	Deployed in research; moving toward commercial [EST]
Atomic clocks	Trapped ions or atoms	GPS, telecommunications, financial timestamping	Global infrastructure [EST]
Rydberg atom sensors	Highly excited atoms	Microwave and radio-frequency field sensing	Rapidly advancing research [EST]

Quantum sensing is not speculative. It is deployed. It generates revenue. It saves lives (MEG brain imaging guides neurosurgery). And it does not require solving error correction. [EST]

Post-Quantum Cryptography — The Urgent Priority

The most immediate practical consequence of quantum computing research is not what quantum computers can do but what they **might** do: break the public-key cryptography that secures the internet. [EST]

The timeline for this threat — often called “Q-Day” — is uncertain. Estimates range from 2030 to 2050 or beyond, depending on assumptions about error correction progress. But here is the key insight: **you do not need to know exactly when Q-Day arrives to need to prepare for it.** Data encrypted today can be stored and decrypted later, once quantum computers become available. This is the “harvest now, decrypt later” threat. If your data must remain secure for 10, 20, or 30 years, you need post-quantum security now — not when Q-Day arrives. [EST]

The good news: the National Institute of Standards and Technology (NIST) finalized post-quantum cryptography standards in 2024. These are classical cryptographic algorithms — running on ordinary computers — that are designed to resist attacks by quantum computers. The standards include:

Algorithm	Type	Purpose
CRYSTALS-Kyber	Lattice-based key encapsulation	Replacing RSA/ECC for key exchange
CRYSTALS-Dilithium	Lattice-based signatures	Replacing RSA/ECDSA for digital signatures
FALCON	Lattice-based signatures	Alternative to Dilithium
SPHINCS+	Hash-based signatures	Stateless, conservative fallback

[EST]

The migration to PQC is underway. Major cloud providers, browser vendors, and operating system developers are integrating these algorithms. The message for organizations in 2026 is clear: **start your PQC migration now**. Inventory your cryptographic assets. Identify long-lived secrets. Plan the transition. This is not speculative. It is not a research project. It is operational risk management. [EST]

Quantum Annealing — The Quiet Achiever

While gate-model quantum computers (IBM, Google, IonQ) have dominated headlines, D-Wave Systems has been shipping **quantum annealers** — specialized quantum processors that solve optimization problems — for over a decade. [EST]

Quantum annealing is not universal quantum computing. It cannot run Shor’s algorithm. It cannot break cryptography. It solves a narrower class of problems: finding the minimum of a complex energy landscape, which maps to optimization problems in logistics, finance, drug discovery, and materials science. [EST]

In early 2026, D-Wave reported the first demonstration of **quantum annealing advantage** on real optimization problems — not contrived benchmarks, but problems with practical relevance. [EST] The advantage was not asymptotic in the sense of Shor’s algorithm. It was an empirical demonstration that, for specific problem instances, the quantum annealer found better solutions faster than classical solvers. This is the kind of practical quantum advantage that gate-model computers have not yet demonstrated. [EST]

D-Wave’s annealers operate at approximately 15 mK — still requiring dilution refrigeration — but they do not require active error correction. The annealing process is relatively robust to noise; the system naturally settles into low-energy states. This makes quantum annealing a different category from gate-model quantum computing — one with a fundamentally different relationship to the thermodynamic wall. [EST]

Quantum Communication — QKD Networks and Satellite Links

Quantum key distribution (QKD) uses quantum states — typically single photons — to distribute cryptographic keys with information-theoretic security. Any eavesdropping attempt disturbs the quantum states and is detectable. [EST]

QKD has been demonstrated over fiber-optic links spanning hundreds of kilometers and via satellite links between ground stations and orbiting spacecraft (notably, China’s Micius satellite). Several countries are building QKD networks for government and financial communications. [EST]

The limitations: QKD requires dedicated optical infrastructure (fiber or line-of-sight to satellites). It does not scale to the internet’s point-to-anywhere architecture. It is best suited for high-security backbone links between fixed nodes. And it competes with PQC, which runs on existing classical infrastructure. The “QKD vs. PQC” debate is active and unresolved — but for most organizations, PQC migration is the practical priority, while QKD is a complementary technology for specific high-security use cases. [EST]

What Does NOT Deliver Today (Despite the Headlines)

- **Factoring large numbers** — no quantum computer has factored a number larger than 21, and that required pre-compiled knowledge of the factors. [EST]
- **Molecular simulation** — no quantum computer has simulated a molecule that a classical computer cannot simulate. Estimates place practical quantum chemistry simulation 5-10 years away, contingent on fault tolerance. [PROP]
- **Quantum machine learning** — no demonstration of quantum advantage for machine learning exists on any real dataset. Theoretical analysis suggests classical ML with kernel methods can often match or exceed quantum approaches. [EST]
- **Breaking RSA** — despite periodic headlines, no quantum computer has broken any real cryptographic system. The threat is real but the capability does not yet exist. [EST]
- **General-purpose quantum computing** — nothing resembling a “quantum CPU” that can run arbitrary programs exists, and the thermodynamic wall suggests it may never exist in the form that popular accounts imagine. [PROP]

Chapter Summary — Part III

Chapter	Key Insight	Confidence
8. The Surface Code Plateau	Being “below threshold” means error correction works — but the gap to useful error rates is 13 orders of magnitude	[EST] for data; [PROP] for extrapolation
9. The Cryogenic Arithmetic	The 20,000× cooling gap is Carnot-limited; at scale, QEC power exceeds feasible cryogenic capacity	[EST] for thermodynamics; [PROP] for QEC extrapolation
10. What Quantum Delivers Today	Quantum sensing delivers now. PQC migration is urgent. D-Wave annealing shows real advantage. Gate-model factoring/simulation/ML do not deliver yet.	[EST]

References — Part III

1. Google Quantum AI. (2023). “Suppressing quantum errors by scaling a surface code logical qubit.” *Nature*, 614, 676-681. [EST]
2. Acharya, R. et al. (2025). “Quantum error correction below the surface code threshold.” *Nature*, 638, 920-926. [EST]
3. Quni-Gudzinis, R. B. (2026). *Bruhat-Tits Quantum Processor*. Zenodo. DOI: 10.5281/zenodo.20109835. [PROP]
4. Quni-Gudzinis, R. B. (2025). *Thermodynamic Constraints and Architectural Inversions*. Zenodo. DOI: 10.5281/zenodo.17938113. [PROP]

5. NIST. (2024). “Post-Quantum Cryptography Standards.” FIPS 203, 204, 205. [EST]
 6. Hempel, C. (2026). “Trapped-Ion Quantum Computers.” In *Quantum Technologies*. Springer. DOI: 10.1007/978-3-031-90727-2_2. [EST]
-

End of Part III. Part IV (Chapters 11-14) will cover The Geometric Alternative: ultrametric geometry, the Bruhat-Tits tree, tree codes, and the threshold advantage — the chapter that makes this guide different from every other quantum computing textbook.

Part IV: The Geometric Alternative

Chapter 11: Geometry Matters — Archimedean vs. Ultrametric

The Hidden Assumption

Every quantum computing textbook — every research paper, every corporate roadmap, every popular article — is built on a mathematical assumption so fundamental that it is almost never stated. That assumption is **Archimedean geometry**. [EST]

Archimedean geometry is the geometry of everyday experience. It is the geometry of Euclidean space, of real numbers, of continuous manifolds. Its defining property is that distances add linearly: if you walk from point x to point y and then from y to z , the total distance is $d(x, y) + d(y, z)$. This seems so obvious that it barely seems worth stating. But it is not the only possible geometry. [EST]

In an Archimedean space, small errors can accumulate. A qubit that suffers a small phase error, then another small phase error, then another, gradually drifts away from its intended state. This is why quantum computers need error correction: in Archimedean geometry, errors accumulate without bound unless actively corrected. [EST]

Ultrametric Geometry — A Different Way to Measure Distance

There exists another geometry — **ultrametric geometry** — where distances do not add in this way. In an ultrametric space, the **strong triangle inequality** holds. (The ordinary triangle inequality says the direct path is never longer than the indirect one: $d(x, z) \leq d(x, y) + d(y, z)$. The strong version is stricter:

$$d(x, z) \leq \max(d(x, y), d(y, z))$$

This inequality is stronger than the ordinary triangle inequality ($d(x, z) \leq d(x, y) + d(y, z)$). Its consequence is profound: **all triangles are isosceles with the two equal sides at least as long as the third**. Small distances cannot accumulate into large ones. [EST]

In an ultrametric space, two “balls” (sets of points within a fixed distance of a center) are either **nested** (one entirely inside the other) or **disjoint** (they do not overlap at all). They never partially overlap, the way two circles drawn on paper

might. This “nested-or-disjoint” property is the geometric source of ultrametric error confinement. [EST]

Think of a family tree. You and your sibling share a parent — you are “close” in the tree (distance 1). You and your first cousin share a grandparent — you are “further apart” (distance 2). You and a stranger share only a distant ancestor — you are “very far apart” (large distance). The family tree is an ultrametric space: everyone at the same generational distance from you is equally distant, regardless of which branch they belong to. This is fundamentally different from Euclidean distance, where two people standing side by side are close regardless of ancestry. [EST]

Why This Matters for Quantum Computing

In a quantum computer built on Archimedean geometry — every existing quantum computer — physical qubits are arranged in a grid or a chain. Errors on one qubit can propagate to neighboring qubits. Correcting those errors requires active measurement: syndrome qubits, classical decoding, feedback loops. All of this generates heat. [EST]

In a quantum computer built on ultrametric geometry, physical qubits are arranged in a **tree** — a hierarchical structure where qubits at the “leaves” of different branches cannot directly interact. An error on one leaf is confined to its branch. It cannot propagate to a leaf on a different branch. Error correction becomes **passive** — a property of the hardware geometry, not an active measurement protocol. [PROP]

This is not a better error correction code. It is a **different geometry for computing** — one where the mathematical properties of the space itself suppress errors. [PROP]

Chapter 12: The Bruhat-Tits Tree — Geometry Becomes a Tree

p -adic Numbers — An Alternative Way to Measure Distance

The p -adic numbers ($\mathbb{Q}_p$) are a completion of the rational numbers — like the real numbers, but using a different notion of distance. In the real numbers, two numbers are “close” if their difference is small. In the p -adic numbers, two numbers are “close” if their difference is **divisible by a high power of a prime p** . [EST]

For example, in the 2-adic numbers: - **1** and **3** differ by **2**, which is divisible by **2¹**. Distance: **1/2**. - **1** and **5** differ by **4**, which is divisible by **2²**. Distance: **1/4**. - **1** and **17** differ by **16**, which is divisible by **2⁴**. Distance: **1/16**.

The p -adic distance between two integers is **smaller** when their difference is divisible by a **higher** power of p . This is the opposite of ordinary distance — but it is a perfectly consistent, mathematically rigorous metric. It satisfies all the axioms of a distance function. And it is **ultrametric**. [EST]

The Bruhat-Tits Tree — The Geometry of p -adic Space

The p -adic numbers have a beautiful geometric realization: the **Bruhat-Tits tree** $\mathcal{T}_p$. This is an infinite **tree** — a graph with no cycles — where each vertex has exactly $p + 1$ neighbors. The tree is $(p + 1)$ -regular. For the simplest case, $p = 2$, the tree is 3-regular: each vertex connects to exactly three others. [EST]

Vertices of the tree correspond to **homothety classes of lattices** in $\mathbb{Q}_p^2$ — equivalence classes of grids in a two-dimensional p -adic vector space. The “boundary” of the tree — the set of all infinite paths starting from a fixed root — is the projective line $\mathbb{P}^1(\mathbb{Q}_p)$. It is a purely discrete, combinatorial object. [EST]

The ultrametric distance between two boundary points x and y is:

$$d(x, y) = p^{-\text{depth}(\text{lca}(x, y))}$$

where $\text{lca}(x, y)$ is the **lowest common ancestor** — the deepest vertex that lies on the paths from the root to both x and y . The deeper the common ancestor, the closer the points. [EST]

This is exactly the family tree analogy from Chapter 11. You and your sibling share a parent — depth 1. You and your cousin share a grandparent — depth 2. The Bruhat-Tits tree makes the “family tree” structure mathematically precise. [EST]

Why This Tree Can Be Hardware

The Bruhat-Tits tree is not just a mathematical abstraction. It has three properties that make it directly implementable as a quantum hardware topology:

1. **It is discrete.** Unlike continuous manifolds, the tree is a graph — vertices and edges. This means it can be built from discrete physical components (qubits arranged in a physical tree structure, with connections only between parent and child nodes).
2. **It is hierarchical.** The tree has a natural root-to-leaf structure. This maps naturally to modular hardware architectures, where “coarse” qubits (near the root) encode logical information and “fine” qubits (at the leaves) provide redundancy.
3. **It supports local operations.** Error correction on the tree requires only parent-child and sibling interactions — no long-range connectivity. Each subtree can be decoded independently, in parallel. [PROP]

The Bruhat-Tits Quantum Processor (BTQP) paper proposes a concrete hardware topology: a truncated tree of depth d (typically 5-11), with leaf vertices serving as physical qubits and internal vertices implementing majority-vote error correction. [PROP]

Chapter 13: The Tree Code — Error Correction Without Measurement

The Basic Building Block — The $[[3, 1, 1]]$ Perfect Tensor

The tree code is built from a simple quantum error-correcting code at each internal vertex of the tree. The $[[3, 1, 1]]$ code encodes 1 logical qubit into 3 physical qubits and can correct a single arbitrary error. [EST]

The notation $[[n, k, d]]$ means: n physical qubits encode k logical qubits with code distance d (the minimum number of physical errors that cause a logical error). For the $[[3, 1, 1]]$ code: 3 physical qubits, 1 logical qubit, distance 1 (can detect 1 error but cannot fully correct an arbitrary error — it can correct bit-flip OR phase-flip, not both simultaneously without concatenation). [EST]

Concatenation — Building a Hierarchy of Protection

The power of the tree code comes from **concatenation**: placing $[[3, 1, 1]]$ codes at every internal vertex of the Bruhat-Tits tree and nesting them. A physical qubit at a leaf feeds into its parent vertex’s code. That parent’s logical output feeds into the grandparent’s code. And so on, up to the root. [PROP]

At each level, errors are corrected locally — within that vertex’s three qubits — before the logical information is passed upward. An error at a leaf is corrected at its parent. An error at a subtree is corrected at that subtree’s root. The hierarchy of the tree provides a hierarchy of error protection. [PROP]

Holographic Encoding — Bulk and Boundary

The tree code uses a **holographic** encoding: logical information lives at the “bulk” (internal vertices, near the root), while physical qubits live at the “boundary” (leaf vertices). This is directly analogous to the AdS/CFT correspondence in theoretical physics, where a gravitational theory in a bulk spacetime is equivalent to a quantum field theory on its boundary. [PROP]

In the tree code, the holographic principle has a concrete, implementable form: the logical state at the root is redundantly encoded across all leaves, with the tree structure providing the error-correction mechanism. Errors at the boundary are “washed out” as information propagates inward — exactly as in the AdS/CFT correspondence, where boundary perturbations are geometrically suppressed in the bulk. [PROP]

Why This Is Passive Fault Tolerance

The key difference between the tree code and the surface code:

Aspect	Surface Code	Tree Code
Error correction mechanism	Active syndrome measurement + classical decoding	Geometric confinement by tree structure
Syndrome measurements per cycle	120 (for $d = 11$)	0 [PROP]
Decoding	Global MWPM ($O(N^3)$)	Local, parallel ($O(s \log s)$) [PROP]
Encoding rate	$k/n \rightarrow 0$ (vanishes with distance)	$1 - 1/p$ (constant; $\sim 1/2$) [PROP]
Physical qubits (distance-11 equivalent)	241	~ 40 [PROP]

Aspect	Surface Code	Tree Code
Operating temperature	~ 10 mK (dilution refrigerator)	~ 4 K (pulse-tube cryocooler) [PROP]

The tree code does not eliminate errors — no physical system can. What it does is geometrically **confine** errors to their local branches, preventing them from accumulating across the system. This is passive fault tolerance: error suppression as a property of hardware geometry, not active measurement protocols. [PROP]

Chapter 14: The Threshold Advantage — By the Numbers

Where These Numbers Come From

The thresholds quoted below are from the Bruhat-Tits Quantum Processor paper (Quni-Gudzinās, 2026, DOI: 10.5281/zenodo.20109835). They have been verified by three independent methods: **analytical recursion** (mathematical derivation of threshold conditions), **exact Gottesman-Knill stabilizer simulation** (a classical algorithm that efficiently simulates certain classes of quantum circuits), and **Monte Carlo methods** (randomized numerical simulation). [PROP]

Important caveat: These thresholds have been verified by classical simulation — not by physical quantum hardware. The tree code has been computationally validated but not experimentally demonstrated on physical qubits. This is an acknowledged gap. The thresholds are [PROP] — framework claims supported by classical evidence — not [EST] — experimentally confirmed facts. This distinction matters. [GAP]

Bit-Flip Errors — 4.6× Better

The bit-flip channel models errors where a qubit’s state flips from $|0\rangle$ to $|1\rangle$ (or vice versa) with probability p . This is the simplest error model.

Code	Threshold p_c
Surface code (MWPM decoding)	~ 10.9% [EST]
Tree code (BTQP)	50.0% [PROP]
Advantage	4.6×

Interpretation: a physical qubit platform with a 10% bit-flip rate is above the surface code threshold (error correction fails) but far below the tree code threshold (error correction succeeds with room to spare). [PROP]

Depolarizing Errors — 75× Better

The depolarizing channel is more realistic: with probability p , the qubit undergoes a completely random Pauli error (X , Y , or Z with equal probability). This is the standard benchmark for QEC performance.

Code	Threshold p_c
Surface code (MWPM decoding)	$\sim 1.0\%$ [EST]
Tree code (BTQP)	75.0% [PROP]
Advantage	75×

Interpretation: the surface code requires physical qubits with better than 99% fidelity to function. The tree code tolerates qubits that fail three times out of four. This is not an incremental improvement — it is a qualitative change in the physical requirements for fault-tolerant quantum computing. [PROP]

Independent $X + Z$ Errors — A New Benchmark

Code	Threshold p_c
Tree code (BTQP)	17.30% [PROP]

The independent $X + Z$ channel models bit-flip and phase-flip errors occurring independently. The surface code does not have a directly comparable threshold for this channel (its performance on X and Z errors is inherently asymmetric). The 17.30% threshold establishes a new benchmark. [PROP]

Beyond Thresholds — Resource Scaling

Threshold comparisons capture one dimension of code performance. Resource scaling captures another — and often more important — dimension:

Metric	Surface Code	Tree Code
Encoding rate (fraction of qubits doing useful computation)	$k/n \rightarrow 0$ as distance increases	$1 - 1/p$ (constant; $\sim 1/2$ for $p = 2$) [PROP]
Physical qubits for distance-11 equivalent protection	241	~ 40 [PROP]
Decoding time	$O(N^3)$ (global matching)	$O(s \log s)$ (local, parallel) [PROP]
Syndrome measurement overhead	120 measurements per cycle	0 (passive) [PROP]
Operating temperature	~ 10 mK	~ 4 K (200–400× warmer) [PROP]

What the Validation Shows

The computational validation (DOI: 10.5281/zenodo.20134944) tested the tree code under controlled noise at varying tree depths:

- **Zero logical errors** in 500 trials at physical error rate $p_{\text{err}} = 0.40$ for tree depths $d \geq 3$. Equivalent flat encodings failed with logical error rates up to 0.152. [PROP]
- **Energy barrier scaling:** $E_{\text{barrier}}(d) = 2^d$, confirmed exhaustively for $d = 2, 3$ and analytically to $d = 10$. [PROP]
- **Strong triangle inequality:** Verified with zero violations in 15,000 random trials across primes $p = 2, 3, 5$. [PROP]

What This Means for the Thermodynamic Wall

Recall from Chapter 9: the surface code faces a thermodynamic wall because active error correction — syndrome measurements, classical decoding, feedback loops — generates heat that overwhelms cryogenic cooling at scale. The Carnot limit imposes a 20,000× gap between cooling at 4 K (1 W) and at 20 mK (50 μ W). [EST]

The tree code operates at 4 K — **200× warmer** than the surface code’s 20 mK — on standard pulse-tube cryocoolers. No dilution refrigeration is required. The cooling budget at 4 K is 1 watt — 20,000× larger than at 20 mK. This difference is not incremental. It is the difference between a physically feasible cooling budget and one that violates thermodynamic constraints at scale. [PROP]

This is the ultimate significance of the threshold advantage: it is not just that the tree code tolerates more errors. It is that the tree code’s passive architecture — zero syndrome measurements, local decoding, constant encoding rate — operates within thermodynamic constraints that the surface code’s active architecture cannot satisfy at commercially useful scales. [PROP]

Chapter Summary — Part IV

Chapter	Key Insight	Confidence
11. Geometry Matters	Archimedean geometry (errors accumulate) is an assumption, not a requirement. Ultrametric geometry (errors confined) is a legitimate alternative	[EST] for mathematics; [PROP] for QC application
12. The Bruhat-Tits Tree	$\mathcal{T}_p$ is a $(p + 1)$ -regular tree that geometrically realizes p -adic ultrametric space — discrete, hierarchical, directly implementable as hardware topology	[EST] for mathematics; [PROP] for hardware realization
13. The Tree Code	Hierarchical concatenation of $[[3, 1, 1]]$ perfect tensors enables passive fault tolerance: zero syndrome measurements, $O(s \log s)$ decoding, constant encoding rate	[PROP]

Chapter	Key Insight	Confidence
14. The Threshold Advantage	75× depolarizing advantage over surface codes; 4.6× bit-flip; zero logical errors demonstrated at $p = 0.40$; 4 K operation avoids the thermodynamic wall	[PROP] — classically validated, not yet experimentally demonstrated

References — Part IV

1. Quni-Gudzinis, R. B. (2026). *Bruhat-Tits Quantum Processor*. Zenodo. DOI: 10.5281/zenodo.20109835. [PROP]
2. Quni-Gudzinis, R. B. (2026). *Ultrametric Quantum Computation — An MVP Program*. Zenodo. DOI: 10.5281/zenodo.20014913. [PROP]
3. Quni-Gudzinis, R. B. (2026). *Computational Validation of Ultrametric Error Confinement*. Zenodo. DOI: 10.5281/zenodo.20134944. [PROP]
4. Quni-Gudzinis, R. B. (2026). *Ultrametric Quantum Computing: Foundations, Evidence, and Falsifiable Predictions*. QNFO/.github releases. [PROP]
5. Quni-Gudzinis, R. B. (2026). *A Different Geometry for Computing*. QNFO/.github releases. [PROP]
6. Gubser, S. S., Knaute, J., Parikh, S., Samberg, A., & Witaszczyk, P. (2016). “p-adic AdS/CFT.” arXiv:1605.01061. [EST]
7. Marcolli, M. (2018/2020). “Holographic Codes on Bruhat-Tits Buildings.” arXiv:1801.09623 / PAMQ 16(1), 1-33. [EST]
8. Rammal, R., Toulouse, G., & Virasoro, M. A. (1986). “Ultrametricity for physicists.” *Reviews of Modern Physics*, 58(3), 765-788. [EST]
9. Quni-Gudzinis, R. B. (2026). *Ultrametric Foundation: A Mathematical Thesis on Non-Archimedean Physics*. Zenodo. DOI: 10.5281/zenodo.21991899. [PROP]
10. Quni-Gudzinis, R. B. (2026). *The Ultrametric Paradigm*. Zenodo. DOI: 10.5281/zenodo.21992229. [PROP]
11. Quni-Gudzinis, R. B. (2026). *Quantum Laws of Form*. Zenodo. DOI: 10.5281/zenodo.21991953. [PROP]

End of Part IV. Part V (Chapters 15-17) will cover The Falsifiable Science: confidence tagging as a methodology, the three experiments (E1-E3) that can test the ultrametric framework, and scenarios for what happens if they pass or fail.

Part V: The Falsifiable Science

Chapter 15: Why Falsifiability Matters

The Problem with Quantum Computing Roadmaps

For twenty years, the quantum computing industry has issued roadmaps promising fault-tolerant quantum computers in “five to ten years.” These roadmaps are specific about what will be built — qubit counts, gate fidelities, code distances — but they are rarely specific about what would count as **failure**. If a roadmap says “1,000

logical qubits by 2030” and 2030 arrives with 100 logical qubits, the roadmap is simply updated to “1,000 logical qubits by 2035.” The goalposts move, and the promise remains “five to ten years away.” [EST]

This is not a criticism of any specific company. It is a structural feature of how the industry communicates. Corporate roadmaps are marketing documents as much as engineering plans. They are designed to maintain investor confidence, attract talent, and sustain enthusiasm. They are not designed to be falsifiable. [EST]

This guide has presented an alternative: the ultrametric quantum computing framework, which makes specific predictions that **can be wrong**. This chapter explains why falsifiability matters — not just for this framework, but as a general principle for evaluating any scientific or technological claim.

What “Falsifiable” Means

A claim is **falsifiable** if there exists a possible observation that would prove it wrong. The statement “all swans are white” is falsifiable because observing a single black swan would disprove it. The statement “quantum computers will be commercially useful within ten years” is less falsifiable because “commercially useful” is vague and “within ten years” slides forward each year. [EST]

The philosopher Karl Popper argued that falsifiability is the demarcation criterion between science and non-science. A theory that can explain any possible observation — that is compatible with every conceivable outcome — explains nothing. It is unfalsifiable. [EST]

The confidence tagging system used throughout this guide — [EST], [PROP], [GAP], [SPEC], [OPEN] — is a practical implementation of falsifiability. Every [PROP] tag marks a claim that is logically consistent and supported by available evidence but has not yet survived a definitive test. The tag is an invitation to test it. [PROP]

The Self-Scoring Methodology

The QWAV research program includes a self-assessment document — the “Honest Investment Assessment” — that scores the program’s readiness across ten criteria on a 1-10 scale. The overall score: **5.6 out of 10**. [EST]

The lowest scores went to computational validation (2/10) and team (2/10) — precisely the areas where the program is weakest. The highest scores went to problem validity (9/10), solution novelty (8/10), and timing (8/10) — the areas where the program is strongest. This self-scoring is published openly. It invites contradiction. If you think the problem validity should be 4/10 instead of 9/10, the evidence is available for you to make that case. [EST]

This is the methodology this guide recommends: **score your claims. Publish your scores. Invite contradiction**. It costs nothing. It builds credibility with audiences that evaluate on substance. And it forces you to be honest with yourself about what you know and what you hope. [EST]

The program’s falsification discipline has at least one worked example on the record. A 2026 follow-on to the program’s number-theory work (10.5281/zenodo.21979060) subjected the branch-depth reading of stabilizer codes to its own stated falsifiability condition — and published a correction, not a confirmation: the naive mapping

carries no new content, and code distance does not admit a valuation reading. A research program that publishes its own refutations is, at minimum, running the experiment this chapter recommends. Whether the discipline generalizes is for the reader to judge. [PROP]

What Fair Criticism Looks Like

The ultrametric framework deserves criticism. Every scientific framework does. But criticism is only useful if it engages with the actual claims. Here is what fair criticism of the ultrametric framework would look like — and what it would not:

Fair criticism (engages with the claims): - “The tree code thresholds were computed classically. The claim that they hold for quantum systems is [GAP] — and a significant one. Here is why I think classical simulation may not capture the relevant physics...” - “The thermodynamic wall argument assumes surface-code-specific resource scaling. But LDPC codes may reduce the overhead by a factor of $10\times$. Here is the math...” - “The CMB prediction (E1) assumes a specific coupling between Planck-scale geometry and large-scale structure. Here is why that coupling may not exist...”

Not fair criticism (does not engage with the claims): - “This is not peer-reviewed.” (The work is open-access by design. Peer review is not sought. The credibility mechanism is public availability + falsifiable predictions + computational validation.) - “The author does not have a PhD.” (The work stands or falls on its content, not the author’s credentials.) - “This is too radical.” (Radical $\neq$ wrong. Newton was radical. Einstein was radical.)

The principle: **engage with the claims, not the claimant.** If you believe the BTQP thresholds are wrong, show where the math fails. If you believe the thermodynamic wall is surmountable, show the engineering path. This is how science advances. [EST]

Chapter 16: The Three Experiments

The ultrametric quantum computing framework makes three predictions that can be wrong. These are designated E1, E2, and E3 in the QWAV MVP program document (DOI: 10.5281/zenodo.20014913). If any one prediction is definitively falsified, the framework is wrong. If all three are confirmed, the framework becomes difficult to dismiss. [PROP]

E1: CMB Log-Periodic Oscillations

What it tests: If the early universe has an ultrametric structure at the Planck scale — specifically, if spacetime is fundamentally p -adic rather than Archimedean — then the cosmic microwave background (CMB) power spectrum should exhibit **log-periodic oscillations**: variations in intensity that are periodic in the logarithm of the angular scale ℓ , with period proportional to $\log_p(\ell)$. [PROP]

How to test it: The Planck satellite’s 2018 data release provides the most precise CMB power spectrum ever measured, covering angular scales from $\ell = 2$ to $\ell \approx 2500$. The data is publicly available. The analysis involves searching for

periodic structure in $\log(\ell)$ space using standard spectral analysis techniques (Lomb-Scargle periodogram, wavelet analysis). [EST]

What a positive result means: Log-periodic oscillations in the CMB at predicted frequencies would be evidence for discrete scale invariance at the Planck scale — a signature of ultrametric geometry in the early universe. This would be independent confirmation that ultrametric structure is physically real, not just mathematically elegant. [PROP]

What a negative result means: The absence of oscillations at the predicted frequencies would constrain the coupling between Planck-scale geometry and observable cosmology. It would not falsify the entire framework — the coupling might be too weak to observe, or the specific prediction might need refinement — but it would eliminate the simplest and most accessible test. [PROP]

Parameter	Value
Cost	~ \$60,000 (data analyst time + computational resources)
Timeline	3-6 months
Data	Publicly available (Planck 2018)
Hardware needed	None (computational analysis only)
Risk	Low
Status	[OPEN] — not yet executed

E2: Prime-Modulated Qubit Noise

What it tests: If quantum measurement is fundamentally a projection from a p -adic tree — specifically, if the Born rule arises from the Monna map (a mathematical mapping from ultrametric to Archimedean spaces) — then environmental noise in quantum systems should exhibit peaks at frequencies that are modulated by prime numbers. Specifically, the noise spectrum should show structure at frequencies $f_k = f_q \cdot p^{-k}$ for primes p . [PROP]

How to test it: Run noise spectroscopy on existing cloud-accessible quantum hardware (IBM Quantum, IonQ, QuEra). Measure the noise power spectral density across a wide frequency range. Search for periodic structure in $\log(f)$ space correlated with prime numbers. Compare with control: generate random frequency spectra and test whether the prime-modulated structure appears by chance. [PROP]

What a positive result means: Prime-modulated noise would be direct evidence that quantum measurement has an underlying ultrametric structure. This would be observable on **existing hardware** — no new devices need to be built. The implications would be profound: it would mean the standard (Archimedean) interpretation of quantum measurement is incomplete. [PROP]

What a negative result means: The absence of prime-modulated noise at detectable levels would constrain the strength of the ultrametric signal relative to technical noise. It would suggest either that the Monna map coupling is weaker than predicted, or that current hardware noise floors are too high to detect it. [PROP]

Parameter	Value
Cost	~ \$200,000 (cloud computing time + data analysis)
Timeline	6-12 months
Hardware needed	Existing cloud quantum computers (IBM, IonQ, QuEra)
Risk	Medium (dependent on hardware access and noise floor)
Status	[OPEN] — not yet executed

E3: Tree Architecture Gate Threshold Test

What it tests: If quantum logic gates are fundamentally discrete operations on a Bruhat-Tits tree — specifically, if gate operations correspond to tree automorphisms — then a physical qubit array arranged in a tree topology should exhibit **step-function threshold switching** rather than the smooth $\sin^2$ Rabi oscillations observed in standard (Archimedean) qubit arrays. The transition from “gate off” to “gate on” should be sharp — a threshold phenomenon — rather than gradual. [PROP]

How to test it: Build a small physical qubit array — 10-40 qubits — arranged in a tree topology (e.g., branching factor $p = 3$, depth $d = 3$). Implement quantum gates using standard techniques (microwave pulses for superconducting qubits, laser pulses for trapped ions or neutral atoms). Measure the gate fidelity as a function of control parameters. Compare with an identical number of qubits arranged in a standard (linear or grid) topology. The prediction: the tree topology should show threshold switching where the standard topology shows smooth oscillation. [PROP]

What a positive result means: Threshold switching in a tree topology would be direct evidence that tree geometry enables fundamentally different gate physics. This would be the strongest possible confirmation of the ultrametric framework — a hardware demonstration that geometry alone can improve quantum gate performance. [PROP]

What a negative result means: If the tree topology shows no threshold switching — if its gate characteristics are identical to the standard topology — then the core physical claim of the ultrametric framework is falsified. The mathematics of tree codes would remain valid as a classical approximation, but the claim that tree geometry enables **physically different** gate behavior would be disproven. [PROP]

Parameter	Value
Cost	\$0.5M – \$2M (hardware fabrication + testing)
Timeline	18-36 months
Hardware needed	Custom-fabricated tree-topology qubit array

Parameter	Value
Risk	High (requires new hardware that does not exist)
Status	[PRE-EXPERIMENTAL] — not yet executed; a trapped-ion falsifiability register with five pre-registered observables has been published (10.5281/zenodo.22025544)

#Since the May 2026 edition, the hardware path has a design. The trapped-ion ultrametric testbed (10.5281/zenodo.22025544) organizes twenty-one program records into one testable claim, with five pre-registered observables, each carrying a kill-condition and an apparatus already demonstrated in the cited literature. The register also records what the program has already ruled out: generic clock-rest coupling violates ultrametricity at a 29–35 percent rate, while diagonal coupling yields exact ultrametric structure — the very asymmetry the protocol is designed to measure. This does not execute E3. It makes E3 fundable: the design, the observables, and the kill-conditions are on the table for any laboratory that wants to run the experiment. [PROP]

Why This Sequence Matters

The experiments are structured in a deliberate sequence of increasing cost and increasing decisiveness:

1. **E1 (CMB)** is the cheapest and fastest — it uses existing data and requires only computational analysis. A positive result would be intriguing but not decisive. A negative result would be inexpensive to obtain.
2. **E2 (Qubit noise)** requires cloud quantum computing access — more expensive, but still using existing hardware. A positive result would be stronger evidence (directly observable on quantum devices). A negative result would be moderately informative.
3. **E3 (Tree gate)** is the most expensive and most decisive — it requires building new hardware. But if E1 and E2 are positive, the case for funding E3 becomes compelling. If E1 and E2 are negative, the case for funding E3 becomes weak.

This is how scientific programs should be structured: cheap experiments first, expensive experiments only if the cheap ones justify them. [EST]

Chapter 17: Scenario Planning — Four Futures

The three experiments define a branching tree of possible outcomes. This chapter maps the four most important scenarios — and what each would mean for the ultrametric framework, for quantum computing, and for you as a reader trying to navigate this landscape. [PROP]

Scenario A: The Dream — All Three Pass

Outcome: E1 (CMB oscillations) confirms log-periodic structure. E2 (qubit noise) confirms prime-modulated spectra. E3 (tree gate) confirms threshold switching.

What it means for the ultrametric framework: The framework transitions from “mathematically elegant, computationally validated, not yet experimentally demonstrated” to “experimentally confirmed across three independent domains.” The confidence tags would shift: many [PROP] claims would become [EST]. The framework would be difficult to dismiss. [PROP]

What it means for quantum computing: The industry would face a genuine paradigm challenge — not from a competitor with better qubits, but from a different geometry. The surface code and active QEC would not disappear overnight, but the argument that they face a thermodynamic wall would gain experimental support. Investment would flow toward tree-topology quantum processors. [PROP]

What it means for you: If you are a student, this is the moment to learn ultrametric geometry. If you are an investor, this is the moment to fund tree-topology hardware. If you are a researcher, this is the moment to design the next generation of experiments. But if you are a skeptic, this is also the moment to demand the strongest evidence: independent replication, error analysis, falsification attempts. [EST]

Scenario B: The Partial — E1 and E2 Pass, E3 Fails

Outcome: The cosmological and quantum-noise predictions are confirmed, but the tree gate experiment shows no threshold switching — gate behavior on tree topologies is identical to standard topologies.

What it means: The geometric insight — that ultrametric structure is physically real and observable — would be validated. But the engineering claim — that tree geometry enables different gate physics — would be falsified. This would be a mixed result: the framework is partly right, but not in the way that directly enables a new kind of quantum computer. [PROP]

What happens next: The research program would pivot. The confirmed predictions (CMB oscillations, qubit noise) would be pursued as fundamental physics. The falsified prediction (tree gate switching) would constrain the engineering path. New questions would emerge: Does the ultrametric structure operate at a different energy scale? Does it require different physical substrates? Are there experimental confounds in E3 that need to be ruled out? [PROP]

Scenario C: The Mixed — One Passes, Two Fail

Outcome: One experiment confirms its prediction; the other two do not. For example, E2 (qubit noise) shows prime-modulated spectra, but E1 (CMB) shows no oscillations and E3 (tree gate) shows no threshold switching.

What it means: Partial validation of a specific prediction, without confirmation of the broader framework. This is the most likely outcome for any ambitious scientific program: some things work, some don't. The confirmed prediction becomes the

foundation for further work. The falsified predictions constrain the framework's scope. [PROP]

What happens next: Iteration. The framework is revised to account for which predictions worked and which didn't. New predictions are generated. The cycle continues. This is normal science — not a failure, but the process by which frameworks are refined. [EST]

Scenario D: The Humbling — All Three Fail

Outcome: None of the three predictions is confirmed. CMB shows no oscillations. Qubit noise shows no prime-modulated structure. Tree gates show no threshold switching.

What it means: The ultrametric framework — at least in its current form — is falsified. The mathematics of tree codes remains valid as a theoretical construct. The thermodynamic critique of active QEC remains valid on its own terms. But the specific claim that ultrametric geometry manifests in observable physical phenomena — at the scales and with the signatures predicted — would be disproven. [PROP]

What happens next: The framework is either abandoned or fundamentally revised. The honest researcher publishes the negative results. The scientific community gains valuable information: three specific, falsifiable predictions were tested and found wanting. This is how science progresses — not by confirming theories, but by eliminating those that don't survive testing. [EST]

What to Watch For

Regardless of which scenario unfolds, certain patterns would indicate that the framework is being treated fairly — or unfairly:

Signs of honest engagement: - Independent groups attempt to replicate the predictions - Results — positive or negative — are published openly - The framework is revised in response to evidence - Critics engage with specific claims rather than credentials

Signs of resistance regardless of evidence: - The predictions are ignored rather than tested - The framework is dismissed without engagement (“not peer reviewed,” “no PhD”) - Goalposts move: “the CMB oscillations aren't at exactly the right frequency, so the framework is wrong” despite the prediction being approximate - The framework is never allowed to succeed because success criteria keep changing

The reader of this guide is now equipped to watch for both patterns. You know what the predictions are. You know what would confirm them and what would falsify them. You can follow the evidence wherever it leads — which is the only thing any honest scientist or informed citizen can do. [EST]

Chapter Summary — Part V

Chapter	Key Insight	Confidence
15. Why Falsifiability Matters	“Five to ten years” is unfalsifiable; the ultrametric framework makes predictions that CAN be wrong; self-scoring methodology (5.6/10) invites contradiction	[EST] for philosophy of science; [PROP] for specific framework evaluation
16. The Three Experiments	E1 (CMB: \$60K, 3-6 mo), E2 (qubit noise: \$200K, 6-12 mo), E3 (tree gate: \$0.5-2M, 18-36 mo) — increasing cost, increasing decisiveness	[PROP] for predictions; [OPEN] for experimental status
17. Scenario Planning	Four futures mapped: all pass (paradigm shift), partial (geometric insight validated), mixed (iteration), all fail (framework falsified)	[PROP] for scenarios

References — Part V

1. Popper, K. (1959). *The Logic of Scientific Discovery*. Hutchinson. [EST]
2. Quni-Gudzinis, R. B. (2026). *Ultrametric Quantum Computation — An MVP Program*. Zenodo. DOI: 10.5281/zenodo.20014913. Chapters 15-22. [PROP]
3. Quni-Gudzinis, R. B. (2026). *Honest Investment Assessment — The \$100,000 Question*. QWAV Strategy Archive, May 2026. [EST]
4. Planck Collaboration. (2020). “Planck 2018 results. I. Overview and the cosmological legacy of Planck.” *Astronomy & Astrophysics*, 641, A1. [EST]
5. Quni-Gudzinis, R. B. (2026). *The Trapped-Ion Ultrametric Testbed: A Falsifiability Register*. Zenodo. DOI: 10.5281/zenodo.22025544. [PROP]
6. Quni-Gudzinis, R. B. (2026). *Implications for Computing and Quantum Error Correction*. Zenodo. DOI: 10.5281/zenodo.21979060. [PROP]

End of Part V. Part VI (Chapters 18-20) will complete the guide: *The Blueprint — What to Do in 2026, How to Read the Research, and What We Still Don’t Know*.

Part VI: Synthesis and Next Steps

Chapter 18: The Blueprint — What to Do in 2026

The Landscape, Summarized

This guide has covered a lot of ground. Part I explained why no commercially useful quantum computer exists. Part II taught the standard quantum computing curriculum — qubits, entanglement, algorithms, error correction. Part III explained why the standard approach is stalling: the surface code plateau, the cryogenic arithmetic, the gap between “below threshold” and “usefully low.” Part IV introduced an alternative — ultrametric geometry, Bruhat-Tits trees, passive fault tolerance, and threshold advantages up to 75×. Part V explained how to test the alternative with falsifiable experiments.

Now, the practical question: **what should you actually do?**

The No-Regrets Actions

Some actions are valuable regardless of which quantum computing paradigm ultimately succeeds:

1. Migrate to post-quantum cryptography. Start now.

This is the single most concrete, actionable recommendation in this guide. NIST’s PQC standards (CRYSTALS-Kyber, CRYSTALS-Dilithium, FALCON, SPHINCS+) were finalized in 2024. [EST] Major software vendors are integrating them. The “harvest now, decrypt later” threat means that data encrypted today with RSA or ECC can be captured, stored, and decrypted once quantum computers — of ANY architecture — reach sufficient scale.

If your organization handles secrets that must remain secure for 10+ years, begin your PQC migration now. Inventory cryptographic assets. Identify long-lived secrets. Plan the transition. This is operational risk management, not research curiosity. [EST]

2. Deploy quantum sensing where it adds value.

Quantum sensors — SQUIDs, NV centers, atomic magnetometers, atomic clocks — deliver practical value today without error correction. They improve navigation, medical imaging, geological survey, and precision timing. If your organization operates in these domains, evaluate quantum sensing as an upgrade to existing classical sensor infrastructure. This delivers return on investment now, not in ten years. [EST]

3. Experiment with cloud quantum computing — but budget realistically.

IBM Quantum, IonQ, QuEra, and others offer cloud access to real quantum processors. Running circuits on real hardware teaches you things no textbook can: how noise manifests, why error mitigation is hard, what gate fidelities actually mean in practice. This is valuable education. But budget your time and money as you would for any experimental platform — not as an investment in near-term quantum advantage. The NISQ era continues. [EST]

The Conditional Actions

These actions depend on your risk tolerance, resources, and assessment of the evidence:

4. Fund E1 (CMB log-periodic oscillations). Cost: ~\$60,000.

This is the cheapest, fastest test of the ultrametric framework. It uses publicly available data. It requires only computational analysis — no hardware, no lab, no collaborators. A positive result would not prove the framework, but it would justify further investigation. A negative result would cost little. [PROP]

5. Fund E2 (prime-modulated qubit noise). Cost: ~\$200,000.

This requires cloud quantum computing access but no new hardware. If your organization already has quantum computing budget for experimentation, allocating a portion to noise spectroscopy with ultrametric predictions is a low-marginal-cost addition. [PROP]

6. Monitor E3 (tree gate threshold test). Cost: \$0.5M-\$2M. Do not fund yet.

E3 requires building new hardware. The case for funding it depends on E1 and E2 results. If E1 and E2 are positive, E3 becomes a compelling bet — the decisive experiment. If E1 and E2 are negative, E3 becomes much harder to justify. Wait for data. [PROP]

The Strategic Posture

Platform strategy: Access multiple quantum computing platforms via cloud. Do not bet on a single modality. The cost of platform diversity — learning multiple SDKs, porting algorithms — is far lower than the cost of betting on the wrong hardware. [PROP]

Architecture bet: The highest-leverage research bet in quantum computing is not on any specific qubit platform but on **alternative error correction architectures**. The surface code is hitting thermodynamic limits. The tree code offers 75× better thresholds. LDPC codes, bosonic codes, and other alternatives are being explored. The field is converging on the view that error correction — not qubit count, not gate fidelity — is the bottleneck. [PROP]

Benchmark discipline: demand joules per solution. For any quantum investment — internal or external — require an energy number: total wall energy per correct, useful answer, with cooling and error-correction overhead included. A roadmap without an energy number is a marketing document (Chapter 15). The QWAV energy-efficiency benchmark (10.5281/zenodo.21880104) is one published instance; the metric itself is platform-neutral and costs nothing to demand. [PROP]

Intellectual posture: Distinguish [EST] from [PROP]. Demand falsifiable predictions. Ignore credentials; evaluate evidence. The most valuable skill this guide teaches is not any specific quantum algorithm. It is the ability to read a quantum computing claim and ask: “What would prove this wrong?” [EST]

Chapter 19: How to Read the Research

The Landscape of Quantum Computing Information

Quantum computing information comes from multiple sources, with very different levels of reliability:

Source Type	Examples	Reliability	What to Watch For
Peer-reviewed journals	<i>Nature, Physical Review Letters, Quantum</i>	High for established results; mixed for speculative claims	Peer review catches errors but not hype. “Below threshold” is a mathematical line, not a guarantee of utility.
Preprint servers	arXiv, Zenodo	Variable; depends on author and methods	Open-access by design. Evaluate on content, not venue. Look for falsifiable predictions.
Corporate press releases	IBM, Google, Microsoft blogs	Low for independent assessment; high for corporate intent	Roadmaps are marketing. “Demonstrated X” $\neq$ “X is commercially useful.”
Industry analysts	Gartner, Forrester, McKinsey	Moderate; depends on methodology	Analyst reports reflect consensus but may lag paradigm shifts.
Popular media	<i>Wired, New York Times, Forbes</i>	Low for technical accuracy; high for narrative trends	Headlines overstate. “Quantum breakthrough” headlines appear monthly.
Open-access research	QWAV papers, independent researchers	Variable; evaluate on falsifiability and validation	The open-access model trades peer review for speed and accessibility. Judge on evidence, not credentials.

[EST]

How to Evaluate a Quantum Computing Claim

When you encounter a claim about quantum computing — in a paper, a press release, a news article, a social media post — apply this checklist:

1. **Is the claim falsifiable?** “We will have 1,000 logical qubits by 2030” is specific but the date can slide. “Tree codes achieve a depolarizing threshold of 75%” is specific and can be tested computationally. Falsifiability is the first filter.
2. **What is the confidence level?** Does the source distinguish established fact from framework proposal from speculation? If everything is presented with equal confidence, treat everything with equal skepticism.

3. **What has actually been demonstrated?** “Quantum computer factors 21” is a demonstration. “Quantum computer will factor 2048-bit RSA keys” is a projection. The gap between these two statements is where most hype lives.
4. **What are the resource requirements?** An algorithm that is theoretically $O(n^3)$ may require 20 million physical qubits and 240 kilowatts of power. An asymptotic speedup is not the same as a practical implementation.
5. **Who benefits from the claim?** A company raising money benefits from optimistic roadmaps. An independent researcher publishing open-access benefits from being right. Consider the incentive structure.
6. **What would prove it wrong?** If the answer is “nothing” — if the claim is compatible with any possible observation — it is unfalsifiable. Treat it as marketing, not science.

[EST]

The Confidence Tagging System — A Transferable Skill

The [EST] / [PROP] / [GAP] / [SPEC] / [OPEN] system used throughout this guide is not specific to quantum computing. It can be applied to any domain where claims range from established fact to speculative conjecture:

- **Climate science:** “CO₂ is a greenhouse gas” [EST]. “Global temperatures will rise by exactly 2.7°C by 2100” [PROP].
- **AI safety:** “Large language models can produce harmful outputs” [EST]. “AI will exceed human intelligence by 2040” [SPEC].
- **Medicine:** “Smoking causes lung cancer” [EST]. “This experimental drug will cure Alzheimer’s” [PROP].

The habit of mentally attaching confidence tags to claims — of asking “is this established, proposed, or speculative?” — is a cognitive skill that pays dividends across every domain of knowledge. It is, in the end, the most important thing this guide teaches. [EST]

Chapter 20: What We Still Don’t Know

The Open Problems

This guide has presented the ultrametric quantum computing framework in some detail. It is important to end by acknowledging what the framework does NOT yet answer. Every scientific program has open problems. Listing them honestly is not weakness — it is the map of future work.

1. Perfect tensor existence for $p > 2$. The tree code construction relies on the existence of perfect tensors — mathematical objects that implement the $[[3, 1, 1]]$ code at each tree vertex. Perfect tensors are known to exist for small dimensions, but their existence for arbitrary tree degrees (corresponding to primes $p > 2$) is not proven. This is an open problem in mathematics. [OPEN]

2. Tree automorphism gate generation. The BTQP architecture proposes using tree automorphisms — permutations of the tree that preserve its structure — as logical quantum gates. The set of gates that can be generated this way has not been fully characterized. Whether this gate set is universal — whether it can approximate any quantum computation — is not yet established. [OPEN]

3. Full quantum validation. The tree code thresholds (50%, 75%, 17.30%) were verified by classical simulation — Gottesman-Knill stabilizer simulation and Monte Carlo methods. Classical simulation captures some but not all quantum effects. Full quantum validation — simulating the tree code with superposition, entanglement, and arbitrary noise channels — requires quantum simulation capabilities that exceed current classical computing resources. [GAP]

4. Physical noise modeling. The threshold analysis used idealized error models (independent bit-flip, depolarizing, independent $X + Z$). Real quantum hardware exhibits correlated errors, $1/f$ noise, leakage to non-computational states, and other complex error processes. How the tree code performs under realistic noise models — particularly spatially correlated noise that might violate the assumption of independent subtree errors — is not yet characterized. [GAP]

5. The Adèle ring connection. The p -adic numbers correspond to a single prime p — a single tree. The Adèle ring $\mathbb{A}$ is the mathematical object that combines ALL primes simultaneously — all trees, all at once. The QWAV program contains papers exploring Adelic constraints on quantum field theory and the Langlands program, but the physical interpretation of Adelic quantum computation — what it means to compute on “all trees at once” — remains undeveloped. [SPEC]

6. Hardware realization. No physical tree-topology quantum processor has been built. The five candidate platforms identified in the UQC MVP document (NV centers, neutral atoms, trapped ions, superconducting circuits, twisted superconductors) have not been experimentally tested with tree geometry. The step from mathematical framework to working hardware is large. [GAP]

7. The QEC–Darwinism tradeoff in ultrametric spaces. A 2026 no-go theorem (Maity et al., arXiv:2608.03944) proves that quantum error correction and quantum Darwinism cannot coexist above a critical logical fidelity ($F_L > 0.874$) — in Archimedean geometry. The QWAV program has published an audit of that theorem through Ostrowski’s lens on Bruhat–Tits tree code spaces (10.5281/zenodo.21964674), showing the tradeoff transforms under ultrametric substitution and identifying regimes forbidden by the Archimedean bound at small primes. Whether the transformed bound is experimentally testable on near-term hardware remains open. [OPEN]

8. The calibrated displacement question. When — if ever — will non-Archimedean geometry displace the real numbers as the working substrate of a physical domain? The program’s own structured assessment (10.5281/zenodo.21747228) anchors the question to a historical reference class and registers nine dated calibration predictions. Its central estimate: a 0.15–0.20 probability of displacement in at least one domain by 2036, with the null hypothesis — non-Archimedean modeling remains a niche tool — as the well-anchored prediction. The geometric alternative in this guide is a bet inside that envelope, and the honest reader should price it accordingly. [PROP]

Why Uncertainty Is Not Weakness

A scientific program that acknowledges its open problems is stronger than one that pretends to have all the answers. The open problems are not embarrassments to be hidden. They are the research agenda. They are what a PhD student, a postdoc, or an independent researcher could spend years investigating. They are the invitation to contribute. [EST]

If you have made it to the end of this guide, you understand the landscape well enough to choose an open problem and pursue it. Whether you are a student looking for a thesis topic, an engineer looking for a hardware challenge, or an investor looking for the next bet — the open problems are where the value is. [EST]

Epilogue

This guide began with a question: **why don't we have quantum computers yet?**

The answer, in brief: because the dominant paradigm — active quantum error correction on Euclidean lattices — faces a thermodynamic wall. The $20,000\times$ gap between cooling at 4 K and cooling at 20 mK is a Carnot limit — a hard physical bound, not an engineering optimization problem. The surface code, for all its elegance, cannot cross this wall without violating thermodynamics.

But there is an alternative. Ultrametric geometry — the geometry of hierarchical trees rather than continuous grids — enables **passive** fault tolerance. Error correction becomes a property of the hardware structure, not an active measurement protocol requiring continuous syndrome readout. The tree code achieves thresholds up to $75\times$ higher than the surface code while operating at $200\times$ warmer temperatures, on roughly one-sixth the physical qubit budget (241 vs ~ 40 qubits at distance-11-equivalent protection), with local parallel decoding.

This alternative is published, DOI-registered, and computationally validated. It is not yet experimentally demonstrated. It makes three falsifiable predictions — E1 (CMB oscillations), E2 (qubit noise), E3 (tree gate switching) — that can be tested for under \$300,000 total for E1+E2.

Whether the ultrametric framework proves correct or not, it exemplifies the right way to do science: make specific, falsifiable predictions. Publish open-access. Tag your confidence. Invite contradiction. Let the evidence decide.

The standard quantum computing story will continue. IBM and Google will release new roadmaps. New “quantum supremacy” demonstrations will make headlines. NISQ will persist. The promise will remain “five to ten years away.”

You now know enough to read those headlines critically. You know what has been demonstrated and what hasn't. You know what “below threshold” means — and doesn't mean. You know where the thermodynamic wall is and why it matters. And you know that there is an alternative — not a better surface code, but a different geometry for computing entirely.

What you do with that knowledge is up to you.

Chapter Summary — Part VI

Chapter	Key Insight	Confidence
18. The Blueprint — What to Do in 2026	PQC migration (urgent), quantum sensing (now), cloud experimentation (budget realistically), E1/E2 (fund if possible), E3 (wait for data). Platform diversity is rational. Architecture bet: alternative QEC geometries.	[EST] for PQC/sensing; [PROP] for investment recommendations
19. How to Read the Research	Six-source landscape (journals, preprints, corporate PR, analysts, media, open-access). Six- question evaluation checklist. Confidence tagging as a transferable cognitive skill.	[EST]
20. What We Still Don't Know	Seven open problems: perfect tensors for $p > 2$, tree automorphism gates, full quantum validation, physical noise modeling, Adèle ring connection, hardware realization, and the QEC–Darwinism tradeoff in ultrametric spaces. Uncertainty is the research agenda.	[OPEN] for all seven

References — Part VI

1. NIST. (2024). “Post-Quantum Cryptography Standards.” FIPS 203, 204, 205. [EST]
2. Quni-Gudzinis, R. B. (2026). *Bruhat-Tits Quantum Processor*. Zenodo. DOI: 10.5281/zenodo.20109835. [PROP]
3. Quni-Gudzinis, R. B. (2026). *Ultrametric Quantum Computation — An MVP Program*. Zenodo. DOI: 10.5281/zenodo.20014913. [PROP]
4. Quni-Gudzinis, R. B. (2025). *Thermodynamic Constraints and Architectural Inversions*. Zenodo. DOI: 10.5281/zenodo.17938113. [PROP]
5. Quni-Gudzinis, R. B. (2026). *Convergence, Consilience, and the Hierarchical Architecture of Reality*. Zenodo. DOI: 10.5281/zenodo.20302276. [PROP]
6. Quni-Gudzinis, R. B. (2026). *When Will Non-Archimedean Geometry Displace the Real Numbers?* Zenodo. DOI: 10.5281/zenodo.21747228. [PROP]

End of Part VI. This completes the 20-chapter Revolutionary Beginner's Guide to Quantum Computing. The full guide consists of six parts across six files (0.5.md through 0.10.md), totaling approximately 22,000 words. Supporting documentation

— *landscape analysis (0.1.md, 0.2.md)*, *curriculum blueprint (0.3.md)*, *intellectual genealogy (0.4.md)*, and *table of contents (0.0.md)* — provides additional context and cross-references.

Appendix A: Program Update — August 2026

This appendix records what has changed since the first publication of this guide (May 2026), for readers who already read the earlier version.

What is new in v1.3 (August 2026). This edition adds the joules-per-solution benchmark as a measurable form of the thermodynamic argument (Chapter 9), records the publication of a hardware design for experiment E3 with five pre-registered observables (Chapter 16), documents the program’s published self-correction on prime-valuation readings of quantum codes (Chapter 15), adds the calibrated displacement assessment to the open problems (Chapter 20), and extends the evidence trail with the program’s 2026 thermodynamic-successor analysis and its consolidated mathematical thesis. The five corrections from the May edition, each verified by a deposited computation (Appendix B), remain recorded below.

Corrections. Five quantitative statements in the May 2026 edition were wrong and are corrected in this edition, each verified by an independent computation (Appendix B):

1. The Landauer limit at 20 mK is approximately 1.9×10^{-25} J per erased bit, not 2×10^{-23} J (the earlier figure corresponds to 2 K).
2. The linear extrapolation of Google’s 2023 logical-error data to distance 11 gives approximately 2.6%, not 2.4%.
3. The tree code’s 4 K operating temperature is 200× warmer than 20 mK, not 100×.
4. At distance-11-equivalent protection the tree code uses roughly one-sixth the physical qubit count of the surface code (241 vs ~40), not one-twentieth.
5. A dilution refrigerator runs on a closed helium-3/helium-4 mixture — the 130-ton figure belongs to the LHC, not to a laboratory refrigerator; the LHC’s cryogenic power is likewise corrected to tens of megawatts, not the roughly 40 megawatts stated in the May edition.

The central claims are unchanged: the 20,000× cooling gap, the Carnot arithmetic, and the classically validated tree-code thresholds (50% bit-flip, 75% depolarizing, 17.30% independent X+Z) all stand as published, and all remain tagged [PROP] — computationally validated, not yet demonstrated on quantum hardware.

New evidence records. Since May 2026 the research program has published several records that sharpen or constrain the framework without yet resolving its falsifiable predictions:

- **A falsifiability register for trapped-ion hardware** (10.5281/zenodo.22025544): twenty-one published records organized into one testable claim, with five pre-registered observables, each with a kill-condition and an apparatus already demonstrated in the cited literature. This is the concrete instrumentation path toward experiment E3.
- **The QEC–Darwinism tradeoff in ultrametric spaces** (10.5281/zenodo.21964674): an audit of a 2026 no-go theorem showing the

Archimedean tradeoff transforms on Bruhat–Tits tree code spaces (Chapter 20, open problem 7).

- **A correction on prime-valuation readings of quantum codes** (10.5281/zenodo.21979060): the program’s own falsification discipline produced a correction rather than a confirmation — the naive valuation mapping of stabilizer codes carries no new content, and code distance does not admit a valuation reading. Readers evaluating the program’s honesty should weigh this record alongside its positive claims.
- **The laws-of-form pillar consolidation** (10.5281/zenodo.21991953) and the **ultrametric foundation thesis** (10.5281/zenodo.21991899): the broader consilience program behind the geometric alternative.

Status of the three experiments. E1 (CMB log-periodic oscillations), E2 (prime-modulated qubit noise), and E3 (tree-topology gate threshold test) remain unexecuted as of 2026-08-20. The falsifiability register (above) is the program’s current instrumentation effort toward E3; E1 and E2 still await funding or independent execution. The predictions are unchanged, the kill-conditions are unchanged, and the invitation to test them stands.

Appendix B: Computational Verification

Every quantitative claim in this guide that a computer can check has been checked by the script `artifacts/verification/verify_guide_claims.py`, deposited with this record. The script computes each value independently from first principles — Carnot coefficients of performance, Landauer bounds, code-distance combinatorics, threshold ratios — and asserts agreement with the numbers printed in the text.

Golden values. The key verified quantities:

Quantity	Verified value	Where used
Carnot COP at 20 mK / 300 K	6.67×10^{-5}	Chapter 9
Minimum wall power per watt removed at 20 mK	$\approx 15,000$ W	Chapter 9
Wall power at 30% of Carnot	$\approx 50,000$ W	Chapter 9
Cooling gap (1 W at 4 K vs 50 μ W at 20 mK)	20,000×	Chapters 3, 9, 14
Landauer limit at 20 mK	$\approx 1.9 \times 10^{-25}$ J	Chapter 9
Linear extrapolation to d = 11 (Google 2023 data)	$\approx 2.6\%$	Chapter 8
Errors over 1M gates at that rate	$\approx 26,000$	Chapter 8

Quantity	Verified value	Where used
Surface code $d = 11$ physical qubits ($2d^2 - 1$)	241	Chapters 8, 14
Surface code $d = 11$ stabilizers ($d^2 - 1$)	120	Chapter 14
Bit-flip threshold ratio (50 / 10.9)	$4.6\times$	Chapter 14
Depolarizing threshold ratio (75 / 1.0)	$75\times$	Chapter 14
Tree vs surface qubit count at $d = 11$ (241 / 40)	$\approx 6\times$ (one-sixth)	Chapter 14, Epilogue
4 K vs 20 mK	$200\times$	Chapter 14, Epilogue
Encoding rate $1 - 1/p$ at $p = 2$	$1/2$	Chapter 14
Grover iterations for $N = 10^6$ ($\pi\sqrt{N}/4$)	785	Chapter 6
E1 + E2 combined cost vs \$300,000	\$260,000	Chapters 16–18

The same discipline applies across the program: the QCA Toy Model (10.5281/zenodo.22012694) ships a fully reproducible exact state-vector simulation with the paper, and every reported number in it is reproducible from the deposited source.

Reproducibility. Runtime: Python 3.12, standard library only, deterministic (no random numbers, no seeds required). Re-run with `python artifacts/verification/verify_guide_claims.py`; the script writes `artifacts/verification/verification-log.txt` and exits 0 only if every check passes (35 checks in this edition). The five v1.1→v1.2 corrections are recorded in the log.